

ServiceNow

(CIS-DF)

*Certified Implementation Specialist-Data
Foundations (CMDB and CSDM)*

Valid for 1 Month

Total: 105 Questions

Question: 1

Bilva

Choose 2 option. A Configuration Manager wants to explore ServiceNow CMDB 360 saved queries to see if the reports can assist with managing of CMDB data. What insights are gained from CMDB 360 queries?

- A. Duplicate configuration items from different data sources
- B. Orphan CIs created from different data sources
- C. Gaps in attribute data from different data sources
- D. Different attribute values from different data sources
- E. Unique CIs created from different data sources

Answer: C,D

Explanation:

C. Gaps in attribute data from different data sources

D. Different attribute values from different data sources

Option-Wise Analysis

A. Duplicate configuration items from different data sources (Incorrect)

The identification of duplicates is primarily handled by the Identification and Reconciliation Engine (IRE) using specific Identification Rules. While CMDB 360 uses data processed by the IRE, its specific purpose as a "360" tool is to audit the values within those records, not to act as the primary detector for duplicate CI existence.

B. Orphan CIs created from different data sources (Incorrect)

Orphan CIs are CIs that lack required relationships to other items. This falls under CMDB Health (specifically the Compliance or Correctness pillars). You manage these through the CMDB Health Dashboard and health inclusion rules, rather than the multi-source query engine.

C. Gaps in attribute data from different data sources (Correct)

CMDB 360 includes a specific query type called the "Find Gap" query. This allows a Configuration Manager to see which data sources are failing to provide information for certain CIs. For your project, this is vital for identifying if a source like Qualys is missing from a specific segment of your network.

D. Different attribute values from different data sources (Correct)

This is the core "insight" of CMDB 360. The "Compare Attribute Values" query surfaces conflicts where two or more sources disagree on a value (e.g., one says a server has 8GB of RAM, another says 16GB). This helps you decide which source is more trustworthy and tune your Reconciliation Rules accordingly.

E. Unique CIs created from different data sources (Incorrect)

Unique CIs represent successful identification where a source has found a new asset. While this is part of the overall data flow, it is not considered an "insight" or a problem-solving report that CMDB 360 is designed to generate.

Reference:

<https://www.servicenow.com/docs/r/washingtondc/servicenow-platform/configuration-management-database-cmdb/querying-cmdb.html>

Question: 2

Bilva

Choose 2 option. A ServiceNow Administrator wants to implement Service Graph Connectors to provide integrations to many third-party solutions that the company wants integrated into the CMDB. Which categories of connectors are available to the Administrator?

- A. Workflow Automation
- B. Observability
- C. DevOps
- D. Cloud

Answer: B,D

Explanation:

B. Observability

D. Cloud

Why these are the correct categories:

D. Cloud (Cloud Providers): This is one of the most common categories. It includes major integrations like the **Service Graph Connector for AWS, Azure, and GCP**. These connectors help you pull in virtual machines, storage buckets, and serverless functions directly into the CMDB.

B. Observability (Infrastructure & Monitoring): This category focuses on tools that monitor the health and performance of your stack. Examples include the **Service Graph Connector for Datadog, Dynatrace, and OpenTelemetry**. These provide real-time visibility into how infrastructure is supporting digital services.

Why the others are incorrect:

A. Workflow Automation: While ServiceNow is a workflow automation platform, "Workflow Automation" is not a standard category for the ingestion of CMDB data via Service Graph Connectors.

C. DevOps: Although **Service Graph** itself extends into the DevOps lifecycle (planning and coding), the specific Connectors listed in the store for populating the CMDB are categorized under "Software Development" or specific vendor names rather than a generic "DevOps" category in most administrative consoles.

Question: 3

Bilva

Choose 2 option. A global enterprise integrates data from multiple discovery sources such as ServiceNow Discovery, SCCM, AWS, and manual uploads to populate its CMDB. However, each discovery source categorizes the same CIs differently, leading to duplicate records and inconsistencies across the system. As a result, the CMDB team is struggling with data accuracy and standardization. What actions does the CMDB team take to resolve the issue?

- A. A. Implement and use identification and reconciliation rules to avoid duplicates and standardize CI classification.
- B. B. Allow each discovery source to define its own CI class, even if it results in inconsistencies.
- C. C. Use CI Class Manager to establish standardized CI classes and attributes across all discovery sources.
- D. D. Create a custom script to manually adjust incoming data before storing it in the CMDB.

Answer: A,C

Explanation:

A. Implement and use identification and reconciliation rules to avoid duplicates and standardize CI classification Correct

Identification and reconciliation rules are the core mechanisms in ServiceNow CMDB for maintaining data integrity. Identification rules ensure that each CI from different sources is recognized as the same item, preventing duplicates. Reconciliation rules determine which source is authoritative when there are conflicting data values, ensuring consistent CI classification. By properly configuring these rules, the CMDB team can eliminate duplicate records, standardize CI classifications, and improve overall data accuracy.

B. Allow each discovery source to define its own CI class, even if it results in inconsistencies Incorrect

Allowing each discovery source to define its own CI classes would worsen the problem. This approach leads to more duplicates, inconsistent classifications, and poor data quality. It does not provide any standardization or governance, making the CMDB unreliable.

C. Use CI Class Manager to establish standardized CI classes and attributes across all discovery sources Correct

CI Class Manager in ServiceNow helps the CMDB team define a standardized model for CI classes and their attributes. It allows mapping incoming data from multiple discovery sources to these standardized classes, reducing inconsistencies and improving reporting accuracy. Combining CI Class Manager with identification and reconciliation rules ensures a consistent and reliable CMDB.

D. Create a custom script to manually adjust incoming data before storing it in the CMDB Incorrect

Manually scripting adjustments for incoming data is error-prone, difficult to maintain, and not scalable. ServiceNow provides built-in mechanisms like reconciliation rules and CI Class Manager to handle these tasks automatically and reliably, making custom scripts unnecessary for standard scenarios.

Question: 4

Bilva

A CMDB CI Class Owner responsible for the Windows Servers needs to manage the Windows Server class. Which CI Class Manager feature will help the CI Class Owner streamline this task?

- A. CI Favorites
- B. Pinned Classes
- C. Search CI Classes

Answer: B

Explanation:

B. Pinned Classes

Option-Wise Analysis

A. CI Favorites (Incorrect) While ServiceNow has a global "Favorites" feature (the star icon in the navigation bar), "CI Favorites" is not a specific feature within the **CI Class Manager** interface designed for class owners to streamline the management of the class definitions themselves.

B. Pinned Classes (Correct) This is a specific feature within the **CI Class Manager**. It allows a user to "Pin" the classes they manage or visit most frequently (such as the Windows Server class). Once pinned, these classes appear in a dedicated list, allowing the Class Owner to bypass the entire CMDB hierarchy and jump directly to the configuration settings, health metrics, and identification rules for that specific class.

C. Search CI Classes (Incorrect) Searching is a functional way to find a class, but it is not a "streamlining"

feature in the same way as Pinning. Searching requires the user to manually type the name every time they enter the tool, whereas **Pinned Classes** provides a persistent, one-click shortcut for the Class Owner.

Question: 5

Bilva

A CMDB Administrator is working in the CI Class Manager on the Basic Info tab. How can the class be set as a Principal Class?

- A. Select 'Yes' from the Principal Class choice list
- B. Check the Principal Class check box
- C. Click the Principal Class UI Action button

Answer: B

Explanation:

B. Check the Principal Class check box

Option-Wise Analysis:

A. Select 'Yes' from the Principal Class choice list (Incorrect)

While many settings in ServiceNow use choice lists (dropdowns), the Principal Class designation is a simple toggle. There is no dropdown menu for this specific setting on the Basic Info tab.

B. Check the Principal Class check box (Correct)

In the CI Class Manager, once you have selected a class from the hierarchy and navigated to the **Basic Info** section, you will see a checkbox labeled **Principal Class**. Ticking this box and saving the record designates the class as "Principal," which immediately begins filtering the Configuration Item (CI) lookup fields on task-based records like Incidents and Changes.

C. Click the Principal Class UI Action button (Incorrect)

UI Actions are typically buttons found at the top or bottom of a form (like "Submit" or "Update"). The Principal Class setting is a field value on the form itself, not a separate action button.

Question: 6

Bilva

A CMDB Administrator is using the Duplicate CI Remediator to address a de-duplication task. On the first tab of the wizard, the Main CI is selected. Which attributes are used to identify the Main CI?

- A. Least Related Items
- B. Most Related Items
- C. Newest Created
- D. Oldest Created

Answer: B

Explanation:

A. Least Related Items Incorrect

Selecting a CI with the least related items would not make sense as the Main CI. The Main CI should ideally be

the one with the **most relationships**, so that merging duplicates preserves the maximum amount of data and associations. Choosing the least connected CI risks losing important relationships.

B. Most Related Items Correct

The Duplicate CI Remediator uses the “**Most Related Items**” criterion to identify the Main CI. This ensures that the CI selected as main has the largest number of associated relationships, configurations, and dependencies. Keeping the most connected CI as the main helps preserve the integrity of the CMDB when merging duplicates and prevents loss of critical relationship data.

C. Newest Created Incorrect

Choosing the newest CI is not recommended for the Main CI because newer CIs may have fewer relationships and less historical data. The goal of the Main CI is to retain as much information as possible from all duplicates, which favors the most connected CI rather than the newest one.

D. Oldest Created Incorrect

While the oldest CI may have historical value, it does not necessarily have the most relationships. Duplicate CI Remediator prioritizes **relationship richness** over creation date to maintain CMDB integrity, making “oldest created” an unreliable method for Main CI selection.

Question: 7

Bilva

A CMDB Administrator needs to configure a new application identification rule that considers the potential for the same application installed more than once on the same server. Which is the best choice of a criterion attribute?

- A. Configuration File Name
- B. Version
- C. Class
- D. Configuration File Path
- E. Port

Answer: D

Explanation:

A. Configuration File Name Incorrect

The configuration file name alone is often the same for multiple installations of the same application. If the same application is installed more than once on a server, the file names are usually identical, making this attribute insufficient for uniquely identifying each installation.

B. Version Incorrect

The version attribute helps differentiate between different releases of an application, but it does not uniquely identify multiple installations of the same version on the same server. Two identical installations can share the same version, leading to duplicate or incorrectly matched CIs.

C. Class Incorrect

The CI class simply defines the type of CI, such as Application or Software Instance. It does not provide any uniqueness and cannot distinguish between multiple installations of the same application on a single server.

D. Configuration File Path Correct

The configuration file path uniquely identifies where an application is installed on the server. When the same application is installed multiple times, each installation typically resides in a different directory path. Using the configuration file path as a criterion attribute allows the identification rule to clearly distinguish between separate installations, preventing duplicates and ensuring accurate CI identification.

E. Port Incorrect

Port numbers may help identify running services, but applications can share ports, change ports, or run without fixed ports. Additionally, ports alone do not reliably distinguish multiple installations of the same application on the same server.

Question: 8

Bilva

A CMDB Administrator has installed a Service Graph Connector (SGC), and then made customizations to the mappings. Which is a consequence of this action?

- A. The customized mappings are not supported by ServiceNow, and the customer is responsible for supporting their own customized SGC.
- B. Fields populated by a customization will have a special tag associated with them in the CMDB.
- C. The customization will prevent the SGC from executing without an approval record attached to the affected mapping record.

Answer: A

Explanation:

A. The customized mappings are not supported by ServiceNow, and the customer is responsible for supporting their own customized SGC. Correct

When a Service Graph Connector mapping is customized, it moves outside the standard, out-of-the-box configuration delivered by ServiceNow. As a result, ServiceNow Support does not provide support for issues caused by those customizations. The responsibility for maintaining, troubleshooting, and supporting the customized mappings lies entirely with the customer. This is an important governance consideration when deciding whether to customize SGC mappings.

B. Fields populated by a customization will have a special tag associated with them in the CMDB. Incorrect

ServiceNow does not apply any special tagging or labeling to CMDB fields that are populated through customized Service Graph Connector mappings. Customized and standard fields appear the same in the CMDB, making this option incorrect.

C. The customization will prevent the SGC from executing without an approval record attached to the affected mapping record. Incorrect

Customizing SGC mappings does not introduce any approval requirement. The connector will continue to execute normally without needing approval records. Approval workflows are not enforced as a result of mapping customizations.

Question: 9

Bilva

A CMDB Administrator is leveraging CI data as part of an Integrated Risk Management implementation and the Entity Scoping process. The Administrator wants to leverage the CSDM Data Foundations Dashboard playbooks under the Run tab. Which CSDM relationships are leveraged using the CSDM playbooks?

- A. Business Applications that have relationships to Application Services
- B. Business Applications that have their relationships to Information Objects
- C. Locations that have established parent records
- D. Logical CIs that have relationships with Information Objects

Answer: A

Explanation:

A. Business Applications that have relationships to Application Services Correct

CSDM Data Foundations Dashboard playbooks focus on validating and improving **core CSDM-aligned relationships** that are essential for downstream processes such as Integrated Risk Management and Entity Scoping. One of the most critical relationships is between **Business Applications and Application Services**. These relationships allow ServiceNow to correctly scope risk, assess impact, and associate technical services with business context, which is exactly what IRM depends on.

B. Business Applications that have their relationships to Information Objects Incorrect

While Information Objects are part of the broader CSDM model, the Data Foundations Dashboard playbooks do not focus on validating or enforcing Business Application-to-Information Object relationships. These relationships are not central to entity scoping for IRM.

C. Locations that have established parent records Incorrect

Location hierarchy and parent-child relationships are not a focus of the CSDM Data Foundations Dashboard playbooks. These playbooks concentrate on service and application modeling rather than location structure.

D. Logical CIs that have relationships with Information Objects Incorrect

Logical CI and Information Object relationships are outside the scope of the Data Foundations playbooks. The playbooks prioritize foundational CSDM relationships that directly support service mapping, impact analysis, and risk scoping.

Question: 10

Bilva

The Server [cmdb_ci_server] class uses a dynamic reconciliation rule of lowest value for the Disk Space (GB) field, while the Windows Server [cmdb_ci_win_server] class uses a dynamic reconciliation rule of most reported for the Disk Space (GB) field.Given the following data sources that populate Windows Server data into the CMDB

Data Source	Disk Space (GB)
Tivoli	75
ServiceNow	75
LANDesk	50
Altiris	80

360/Multisource CMDB: Which value would be added to the CMDB for the Disk Space (GB) field of a Windows Server [cmdb_ci_win_server] record?

- A. 80
- B. 75
- C. 50

Answer: B

Explanation:

A. 80 Incorrect

Although 80 GB is the highest value reported, the Windows Server class does not use a “highest value” reconciliation rule. Since this value is reported by only one source, it does not meet the criteria of being the most frequently reported value.

B. 75 Correct

The Windows Server class uses the most reported dynamic reconciliation rule. In this case, 75 GB is reported by two different data sources — Tivoli and ServiceNow — making it the most frequently reported value. As a result, 75 GB is selected and written to the CMDB for the Disk Space field on the Windows Server record.

C. 50 Incorrect

The value 50 GB is reported by only one source. While it is the lowest value, the “lowest value” reconciliation rule applies only to the parent Server class, not the Windows Server class. Since the Windows Server class uses “most reported,” this value is not selected.

Question: 11

Bilva

A Configuration Manager needs to enable a CMDB Data Manager policy to remove records from a CI Class while retaining the ability to restore them within a specified period. Which policy type should the Configuration Manager create?

- A. Archive
- B. Retire
- C. Certification
- D. Delete

Answer: A

Explanation:

A. Archive Correct

An **Archive** policy in CMDB Data Manager removes CI records from active use while storing them in an archive table. This approach preserves the data and allows the records to be **restored within a defined retention period** if needed. Archiving helps keep the CMDB clean and performant without permanently deleting data, making it ideal when restoration may be required.

B. Retire Incorrect

A Retire policy changes the lifecycle state of a CI to retired but does not remove the record from the CMDB. The CI remains in the database and continues to exist for reporting or historical purposes, so it does not meet the requirement to remove records.

C. Certification Incorrect

Certification policies are used to validate and attest to the accuracy of CMDB data. They do not remove or archive records and therefore do not support record restoration requirements.

D. Delete Incorrect

A Delete policy permanently removes CI records from the CMDB. Once deleted, the records cannot be restored, which directly contradicts the requirement to retain the ability to restore them within a specified period.

Question: 12

Bilva

A CMDB Administrator wants to start utilizing the CMDB Health Dashboard and its Key Performance Indicators. What does the Administrator do to start using the dashboard?

- A. Nothing, the dashboard is activated by default
- B. Activate the dashboard scheduled jobs
- C. Download the dashboard from the ServiceNow store
- D. Activate the dashboard system property

Answer: B

Explanation:

A. Nothing, the dashboard is activated by default Incorrect

The CMDB Health Dashboard is available out of the box, but it does not start collecting KPI data automatically. Without enabling the required background processing, the dashboard will not show meaningful health metrics, making this option incorrect.

B. Activate the dashboard scheduled jobs Correct

To begin using the CMDB Health Dashboard, the administrator must activate the **scheduled jobs** that calculate and populate the health KPIs such as Completeness, Correctness, and Compliance. Once these jobs are enabled and run, the dashboard begins displaying meaningful CMDB health data.

C. Download the dashboard from the ServiceNow store Incorrect

The CMDB Health Dashboard is included as part of the CMDB application and does not need to be downloaded separately from the ServiceNow Store.

D. Activate the dashboard system property Incorrect

There is no single system property that activates the CMDB Health Dashboard. The key requirement is enabling the scheduled jobs that generate the KPI data, not toggling a property.

Question: 13

Bilva

The CMDB Administrator group seeks to filter specific CI classes that display on the CMDB Health Dashboard. This ensures that only relevant data is displayed, excluding items that are not ready for management. Which feature can the group utilize to achieve this goal?

- A. Data Refresh Rules
- B. Identification Rules
- C. Reconciliation Rules

Answer: D

Explanation:

A. Data Refresh Rules Incorrect

Data Refresh Rules control how and when CMDB data is refreshed or updated from sources. They do not determine which CI classes appear on the CMDB Health Dashboard, so they cannot be used to filter dashboard visibility.

B. Identification Rules Incorrect

Identification Rules define how CIs are uniquely identified to prevent duplicates. While essential for CMDB accuracy, they do not control inclusion or exclusion of CI classes on the CMDB Health Dashboard.

C. Reconciliation Rules Incorrect

Reconciliation Rules determine which data source wins when there are conflicting attribute values. They have no impact on whether a CI class is included or excluded from CMDB Health reporting.

D. Health Inclusion Rules Correct

Health Inclusion Rules are specifically designed to control which CI classes and records are included in CMDB Health calculations and dashboards. By configuring these rules, the CMDB Administrator group can exclude CI classes that are not ready for management and ensure the dashboard displays only relevant, actionable data.

Question: 14

Bilva

A CMDB Administrator needs the fastest time to value solution for effectively ingesting, managing, and maintaining CIs and relationships. Which management tool will accomplish this to import Windows computer data from SCCM?

- A. IntegrationHub ETL connection to SCCM using Robust Transform Engine (RTE)
- B. Import set using JDBC data source connection to SCCM using transform maps
- C. SCCM Usage Metering Spoke
- D. SCCM Service Graph Connector

Answer: D

Explanation:

A. IntegrationHub ETL connection to SCCM using Robust Transform Engine (RTE) Incorrect

While IntegrationHub ETL with RTE is powerful and flexible, it requires configuration of data sources, transforms, mappings, and ongoing maintenance. This approach increases setup time and complexity, making it unsuitable when the goal is the fastest time to value.

B. Import set using JDBC data source connection to SCCM using transform maps Incorrect

Using JDBC import sets and transform maps is a traditional and highly manual approach. It requires custom mappings, scripting, and continuous maintenance to handle CI relationships and updates, resulting in slower time to value compared to modern CMDB integration methods.

C. SCCM Usage Metering Spoke Incorrect

The SCCM Usage Metering Spoke is designed for collecting software usage and metering data, not for ingesting CMDB CIs and relationships. It does not provide full CI lifecycle management or relationship modeling in the CMDB.

D. SCCM Service Graph Connector Correct

The SCCM Service Graph Connector provides an out-of-the-box, CMDB-aligned integration specifically designed for ingesting Windows computer data from SCCM. It automatically applies Common Service Data Model (CSDM) standards, identification and reconciliation rules, and relationship mapping. This delivers the **fastest time to value** with minimal configuration and ongoing maintenance, making it the best choice for effectively managing CIs and relationships from SCCM.

Question: 15

Bilva

Choose 2 option. An organization is updating the CMDB to include new asset types like IoT devices. Relevant CI classes need to be added and outdated ones need to be removed from the Principal Class filter to ensure accurate display in ITSM processes. Which roles are needed to add or remove classes?

- A. sn_cmdb_admin
- B. personalize_dictionary
- C. cmdb_query_builder
- D. sn_csdm_admin

Answer: A,B

Explanation:

A. sn_cmdb_admin

B. personalize_dictionary

Option-Wise Analysis

A. sn_cmdb_admin (Correct)

This is the core administrative role for the CMDB. It allows the Administrator to access the **CI Class Manager**, navigate the hierarchy, and view the "Basic Info" tab for any class. While it provides the framework to manage classes, it often requires a supporting role to modify specific system-level fields.

B. personalize_dictionary (Correct)

In many ServiceNow instances, the "Principal Class" checkbox is tied to the underlying dictionary definition of the class metadata. Even with sn_cmdb_admin, the field may appear "grayed out" unless the user also has the personalize_dictionary role (or itil_admin). This role allows the user to update the specific attribute that designates a class as Principal.

C. cmdb_query_builder (Incorrect)

This role is strictly for creating and managing saved queries within the **Query Builder**. It does not grant permissions to modify class definitions, metadata, or the Principal Class filter settings.

D. sn_csdm_admin (Incorrect)

While the **CSDM (Common Service Data Model)** provides the strategy for which classes should be Principal

(e.g., Application Services), the sn_csdm_admin role is focused on the service lifecycle and the Data Foundations dashboard. It is not the technical role required to toggle the Principal Class setting in the Class Manager.

Question: 16

Bilva

Choose 2 option. Which ServiceNow solutions automatically create relationships between CI Applications that are part of an Application Service?

- A. IntegrationHub ETL
- B. Service Mapping
- C. Event Management
- D. Discovery
- E. Data Manager

Answer: B,D

Explanation:

A. IntegrationHub ETL Incorrect

IntegrationHub ETL is designed to import and transform data from external systems into ServiceNow. While it can populate CI records, it does not automatically discover or create relationships between applications within an Application Service. Relationship logic must be explicitly defined and maintained.

B. Service Mapping Correct

Service Mapping automatically discovers application components and their dependencies and creates **application-to-application and application-to-infrastructure relationships** within an Application Service. It dynamically maps how applications interact, making it a core solution for building and maintaining accurate Application Service relationships in the CMDB.

C. Event Management Incorrect

Event Management consumes events and alerts and correlates them to existing CIs and services. It relies on CMDB relationships but does not create or discover relationships between CI Applications.

D. Discovery Correct

Discovery automatically identifies applications, software instances, and their relationships on infrastructure components. It contributes relationship data that is used as part of Application Services, especially when combined with Service Mapping. Discovery helps establish foundational CI relationships that support Application Service modeling.

E. Data Manager Incorrect

Data Manager focuses on CMDB data governance, lifecycle management, and cleanup. It does not discover or create CI relationships between applications.

Question: 17

Bilva

A Configuration Manager responsible for a specific region wants to use the CMDB Health Dashboard to improve the data quality of the CMDB for that region. The Configuration Manager only sees the overall score and grouped

by CI Class. How can the Configuration Manager get a score for regionally relevant CIs?

- A. Customize the CMDB Health Dashboard scheduled jobs to group the results by region
- B. Create CMDB groups with type, health, by region
- C. On the CMDB health settings, activate the option, Group scores by region

Answer: B

Explanation:

A. Customize the CMDB Health Dashboard scheduled jobs to group the results by region Incorrect

The CMDB Health scheduled jobs are delivered out of the box and are not designed to be customized for grouping results by attributes such as region. Modifying these jobs would be unsupported and is not the recommended approach for regional scoring.

B. Create CMDB groups with type, health, by region Correct

CMDB Groups allow administrators to define **logical groupings of CIs** based on specific criteria, such as region, location, or ownership. By creating CMDB groups scoped to a particular region and using the **Health** group type, the Configuration Manager can view CMDB Health scores specifically for those regionally relevant CIs. This is the supported and recommended way to obtain regional health insights.

C. On the CMDB health settings, activate the option, Group scores by region Incorrect

There is no CMDB Health setting that allows scores to be grouped directly by region. Regional views are achieved through CMDB Groups, not through a dashboard configuration toggle.

Question: 18

Bilva

Choose 2 option. An organization is using CMDB Query Builder to find all application services with a database that has incidents and all infrastructure in those application services. Which steps does the organization take to build this query?

- A. Add a non-CMDB table to the query
- B. Use a Service Mapping Query to find all incidents related to the database
- C. Use a CMDB Query to include application services and their related infrastructure
- D. Use a Service Mapping Query to include non-CMDB tables like the Incident table

Answer: A,C

Explanation:

A. Add a non-CMDB table to the query Correct

CMDB Query Builder allows you to include **non-CMDB tables**, such as the **Incident** table, in a query. This makes it possible to correlate CMDB CIs like databases or application services with operational data such as incidents. Adding a non-CMDB table is essential to identify databases that currently have incidents associated with them.

B. Use a Service Mapping Query to find all incidents related to the database Incorrect

Service Mapping Queries are used to explore service topology and dependencies within the CMDB. They do not query transactional tables like Incident and are not used to retrieve incident data related to CIs.

C. Use a CMDB Query to include application services and their related infrastructure Correct

A CMDB Query is used to traverse CI relationships and include **application services and all related infrastructure components**, such as servers, databases, and other dependent CIs. This step is required to return the full infrastructure scope for the identified application services.

D. Use a Service Mapping Query to include non-CMDB tables like the Incident table **Incorrect**

Service Mapping Queries cannot include non-CMDB tables. Only CMDB Query Builder supports joining CMDB data with non-CMDB tables like Incident, making this option invalid.

Question: 19

Bilva

Choose 2 option. A CMDB Administrator uses the CMDB Data Foundations Dashboard to gain insights into the CMDB. The results display low scores for several metrics. Which actions can the CMDB Administrator take to improve the CMDB Health?

- A. Focus on metric(s) with Critical and High priorities
- B. Remove non-operational and retired CIs
- C. Adjust the metrics using exclusion rules to improve the scores
- D. Use the Remediation Playbooks linked beside each metric

Answer: A,D

Explanation:

A. Focus on metric(s) with Critical and High priorities **Correct**

The Data Foundations Dashboard assigns priorities to metrics based on their impact on CMDB quality and downstream processes. By addressing **Critical and High priority metrics first**, the administrator targets the most impactful issues affecting CMDB Health. This focused approach delivers meaningful improvements without trying to fix everything at once.

B. Remove non-operational and retired CIs **Incorrect**

While cleaning up retired or non-operational CIs can be beneficial, this action is **not a primary or required step** driven directly from the Data Foundations Dashboard. CMDB Health improvement focuses on fixing data quality issues identified by the metrics rather than broadly removing CIs.

C. Adjust the metrics using exclusion rules to improve the scores **Incorrect**

Exclusion rules should not be used simply to raise health scores. Doing so hides underlying data quality issues instead of resolving them and goes against the intent of the CMDB Health framework.

D. Use the Remediation Playbooks linked beside each metric **Correct**

Each metric in the Data Foundations Dashboard includes a **Remediation Playbook** that provides guided, best-practice steps to correct the underlying data issues. Using these playbooks is a supported and recommended method for improving CMDB Health in a sustainable way.

Question: 20

A CMDB Administrator needs to prevent duplicate CI creation from import sets that load data into the CMDB from vendor shipment files containing CI information. How can the Administrator do this?

- A. Use the CMDBTransformUtil API in the transform script

- B. Create comparison rules in the IRE
- C. Set the system property to utilize the IRE within transform maps
- D. Set the coalesce on two mappings within the transform map

Answer: A

Explanation:

A. Use the CMDBTransformUtil API in the transform script

Detailed Reasoning

The question asks specifically how to prevent duplicates from **Import Sets** loading **vendor shipment files**. In ServiceNow, while the **IRE (Identification and Reconciliation Engine)** is the engine that does the work, standard Transform Maps are not "IRE-aware" by default.

Why A is the Correct Method: To connect a standard Transform Map to the IRE, a developer must use a script. ServiceNow provides the **CMDBTransformUtil** script include specifically for this purpose. Within the onBefore script of a transform map, you call `identifyAndReconcile()`. This passes the import data into the IRE, which then uses its Identification Rules (like Serial Number or MAC Address) to find a match instead of relying on a simple "coalesce."

Why C is a Misconception: While there are system properties that affect the IRE (e.g., `glide.asset.create_ci_with_ire` for assets or `glide.cmdb.transform.enable_ire` for certain Service Graph features), there is **no single system property** that automatically forces all standard legacy Transform Maps to use the IRE. You must explicitly invoke the API via the script described in Option A.

Option-Wise Breakdown

A. Use the CMDBTransformUtil API (Correct):

This is the documented and tested method to apply IRE logic to standard Import Set transforms.

B. Create comparison rules in the IRE (Incorrect):

This is incorrect terminology. The IRE uses **Identification Rules** and **Reconciliation Rules**, but not "comparison rules."

C. Set the system property (Incorrect):

As noted, no global system property enables the IRE for all standard transform maps. This is often a "distractor" answer in exams.

D. Set the coalesce on two mappings (Incorrect):

This is the "legacy" way to prevent duplicates. For the CMDB, this is discouraged because it bypasses the IRE's ability to handle class reclassification and multi-source priority.

Question: 21

A CMDB Data Owner has requested better insights on the different data sources that make up the CMDB data set. The Platform Owner knows that the new Service Graph Connector Central plugin is just what is needed. After installing the plugin, what workspace will have the new Service Graph Connector Central tab available?

- A. CMDB Workspace
- B. Service Graph Connector Workspace

Answer: B

Explanation:

A. CMDB Workspace Incorrect

The CMDB Workspace is focused on CI lifecycle management, health, and governance. While it provides visibility into CMDB data quality and structure, it does not host the Service Graph Connector Central tab, which is specifically designed for managing Service Graph integrations and data sources.

B. Service Graph Connector Workspace Correct

After installing the Service Graph Connector Central plugin, the **Service Graph Connector Central tab appears in the Service Graph Connector Workspace**. This workspace is specifically designed to provide centralized visibility and management of all Service Graph Connectors, including data sources, ingestion status, mappings, and overall CMDB data flow.

C. Discovery Admin Workspace Incorrect

The Discovery Admin Workspace is used to configure and manage Discovery-related activities. Although Discovery data may feed into Service Graph, this workspace does not include the Service Graph Connector Central tab.

Question: 22

Choose 2 option.A CMDB Manager wants to start adding CSDM design and planning (design) domain components into the CMDB.Who is involved in this exercise?

- A. Enterprise Architect
- B. Business Relationship Manager
- C. Application Service Owner
- D. Application Owner

Answer: A,B

Explanation:

A. Enterprise Architect Correct

The Enterprise Architect plays a key role in the **CSDM Design and Planning domain**. This role is responsible for defining enterprise architecture elements such as business capabilities, application portfolios, and high-level application structures. These design-time components are foundational to CSDM and must be modeled correctly in the CMDB before operational data is added.

B. Business Relationship Manager Correct

The Business Relationship Manager represents the business perspective and helps define how business needs map to applications and capabilities. In the design domain, this role contributes to identifying business applications, aligning them with business capabilities, and ensuring the CMDB reflects accurate business context.

C. Application Service Owner Incorrect

The Application Service Owner is primarily involved in the **Manage and Operate domain**, where application

services are running, supported, and monitored. This role is not typically responsible for defining design-time CSDM components.

D. Application Owner Incorrect

The Application Owner focuses on the lifecycle, health, and performance of applications once they are operational. While important in later CSDM stages, this role is not directly involved in the **design and planning domain** activities.

Question: 23

A Configuration Manager working in the CMDB Workspace wants to see how CIs are connected to each other. Which tool can be used?

- A. Business Service Map
- B. Unified Map
- C. Relationship Map

Answer: B

Explanation:

A. Business Service Map Incorrect

The Business Service Map is generated through **Service Mapping** and focuses on visualizing application services and their dependencies. It is not the primary visualization tool used directly within the CMDB Workspace for exploring general CI-to-CI relationships.

B. Unified Map Correct

The **Unified Map** is available within the CMDB Workspace and provides a **modern, interactive visualization** of how CIs are connected to each other. It consolidates relationships from multiple sources and allows Configuration Managers to easily explore, analyze, and understand CI relationships directly from the workspace. This is the recommended tool for viewing CI connectivity in CMDB Workspace.

C. Relationship Map Incorrect

The Relationship Map is the **classic CMDB form-based visualization** used outside of CMDB Workspace. While it does show CI relationships, it is not the primary or recommended mapping tool within the CMDB Workspace experience.

Question: 24

A Configuration Manager needs to ingest third-party CIs into the CMDB. Which method minimizes the risk of technical debt?

- A. Vendor-provided integration
- B. Import Sets and Transform Maps
- C. Table API
- D. Service Graph Connector

Answer: D

Explanation:

A. Vendor-provided integration Incorrect

Vendor-provided integrations can work, but they often vary in quality, CMDB alignment, and long-term support. Many do not fully follow CSDM standards or ServiceNow best practices, which can introduce maintenance challenges and technical debt over time.

B. Import Sets and Transform Maps Incorrect

Import Sets and Transform Maps are flexible but highly customizable. This often leads to complex logic, custom scripts, and ongoing maintenance. Over time, these customizations significantly increase technical debt, especially when CMDB structures or identification rules change.

C. Table API Incorrect

Directly writing data to CMDB tables using the Table API bypasses the Identification and Reconciliation Engine (IRE) and CSDM governance unless carefully scripted. This approach increases the risk of duplicates, inconsistencies, and unsupported configurations, resulting in high technical debt.

D. Service Graph Connector Correct

Service Graph Connectors provide **out-of-the-box, CMDB-aligned integrations** that follow CSDM standards and automatically leverage identification and reconciliation rules. They require minimal customization, are fully supported by ServiceNow, and reduce ongoing maintenance effort. This makes Service Graph Connectors the **lowest-risk and lowest-technical-debt** option for ingesting third-party CI data into the CMDB.

Question: 25

Which shows the most complete list of policy types that are provided by the CMDB Data Manager?

- A. Delete, Attestation, Retire, and Certification
- B. Attestation, Retire, and Certification
- C. Retire, Archive, Attestation, Certification, and Delete
- D. Archive and Delete

Answer: C

Explanation:

A. Delete, Attestation, Retire, and Certification Incorrect

This option is missing key CMDB Data Manager policy types such as **Archive**. Because it does not list all available policy types, it is not the most complete answer.

B. Attestation, Retire, and Certification Incorrect

This list is incomplete and excludes important policy types like **Archive** and **Delete**, making it an incorrect choice.

C. Retire, Archive, Attestation, Certification, and Delete Correct

CMDB Data Manager provides a full set of policy types to manage CI lifecycle and data governance. These include **Retire** for lifecycle state management, **Archive** for removing records while allowing restoration, **Attestation** for data validation, **Certification** for periodic data review, and **Delete** for permanent removal. This

option represents the **most complete and accurate list** of CMDB Data Manager policy types.

D. Archive and Delete Incorrect

Although Archive and Delete are valid policy types, this option does not include all available CMDB Data Manager policies and is therefore incomplete.

Question: 26

Choose 2 option.CMDB Administrator wants to leverage the CMDB Data Foundations Dashboard. What are benefits of using this application?

- A. Uses automation to remediate potential risks
- B. Has a framework to create custom metrics for the CMDB
- C. Checks that important data is valid and properly configured
- D. Provides playbooks to assist in the remediation of potential risks

Answer: C,D

Explanation:

A. Uses automation to remediate potential risks Incorrect

The CMDB Data Foundations Dashboard does not automatically remediate risks. It identifies issues and guides remediation, but the corrective actions must be performed manually or through separate processes.

B. Has a framework to create custom metrics for the CMDB Incorrect

The dashboard provides **out-of-the-box metrics** aligned to CSDM and best practices. It does not offer a framework for creating fully custom CMDB health metrics.

C. Checks that important data is valid and properly configured Correct

One of the primary benefits of the CMDB Data Foundations Dashboard is validating that **foundational CMDB data is complete, correct, and aligned to CSDM standards**. It highlights misconfigurations and missing relationships that can negatively impact ITSM, ITOM, and IRM processes.

D. Provides playbooks to assist in the remediation of potential risks Correct

Each metric in the dashboard includes **Remediation Playbooks** that provide guided, best-practice steps to resolve identified data quality issues. These playbooks help administrators systematically improve CMDB health in a supported and sustainable way.

Question: 27

A CMDB CI Class Owner has been asked to change the icon for the UNIX Server class.Which CI Class Manager tab can the owner use to change the icon for the class?

- A. Suggested Relationships
- B. Basic Info
- C. CI List
- D. Attributes

Answer: B

Explanation:

A. Suggested Relationships Incorrect

The Suggested Relationships tab is used to define or recommend relationships between CI classes. It does not control visual properties such as icons.

B. Basic Info Correct

The **Basic Info** tab in CI Class Manager is where core class properties are configured, including the **icon** used to represent the class in the CMDB and user interfaces. Changing the icon here updates how the UNIX Server class is visually displayed across ServiceNow.

C. CI List Incorrect

The CI List tab displays CI records that belong to the class. It does not provide configuration options for class-level properties like icons.

D. Attributes Incorrect

The Attributes tab is used to manage fields and attributes for the CI class. It does not control the icon or other visual identifiers.

Question: 28

A CMDB team has noticed that some hardware CIs are missing the serial number information, making it difficult to maintain data accuracy. The team needs a structured approach to identify and address these gaps. Which accomplishes this task?

- A. CMDB Data Foundation Playbook
- B. Service Graph Connectors
- C. CI Class Manager

Answer: A

Explanation:

A. CMDB Data Foundation Playbook Correct

CMDB Data Foundation Playbooks are designed to **identify data quality gaps** such as missing mandatory attributes (for example, serial numbers) and provide **guided remediation steps**. This delivers a structured, repeatable approach to improving CMDB data accuracy.

B. Service Graph Connectors Incorrect

Service Graph Connectors focus on **ingesting data from external sources** and mapping it into the CMDB. They do not provide analysis or remediation guidance for existing data quality gaps.

C. CI Class Manager Incorrect

CI Class Manager is used to **configure CI classes**, attributes, and relationships. It does not identify missing attribute values or provide remediation workflows.

Question: 29

A CMDB Administrator wants to leverage dynamic reconciliation rules. Which feature must be enabled?

- A. CMDB 360/Multisource CMDB
- B. CMDB Data Manager
- C. CMDB Workspace
- D. Reconciliation Rules

Answer: A

Explanation:

A. CMDB 360/Multisource CMDB Correct

Dynamic reconciliation rules are a capability of **CMDB 360 / Multisource CMDB**, which allows reconciliation logic to vary by CI class and data source. This feature must be enabled to use dynamic reconciliation behavior.

B. CMDB Data Manager Incorrect

CMDB Data Manager focuses on lifecycle policies such as archive, retire, and delete, not reconciliation logic.

C. CMDB Workspace Incorrect

CMDB Workspace provides visualization and management tools but does not enable reconciliation functionality.

D. Reconciliation Rules Incorrect

Reconciliation rules define behavior, but dynamic reconciliation specifically requires the Multisource CMDB feature to be active.

Question: 30

Where does a user with the appropriate role(s) review and manage the generated tasks after configuring CMDB Data Manager policies?

- A. CMDB Health Dashboard > Audit tab
- B. CMDB Workspace > Management tab
- C. CMDB Workspace > My Work tab
- D. CMDB Health Dashboard > Duplicate CIs tab

Answer: C

Explanation:

A. CMDB Health Dashboard > Audit tab Incorrect

The Audit tab is used to review compliance and historical health-related information. It does not display or manage tasks generated by CMDB Data Manager policies.

B. CMDB Workspace > Management tab Incorrect

The Management tab is used for configuration and administration activities, not for reviewing or working on generated policy tasks.

C. CMDB Workspace > My Work tab Correct

Tasks generated by CMDB Data Manager policies (such as archive, retire, or delete actions) are assigned to users and can be reviewed, tracked, and managed in the **My Work** tab within the CMDB Workspace.

D. CMDB Health Dashboard > Duplicate CIs tab Incorrect

The Duplicate CIs tab focuses on identifying and remediating duplicate configuration items, not managing Data Manager policy tasks.

Question: 31

Choose 2 option. Which are CMDB Data Manager end of life policy types?

- A. Lost
- B. Decommission
- C. Archive
- D. Disposed
- E. Retire

Answer: C,E

Explanation:

A. Lost Incorrect

Lost is an **asset lifecycle status** used in Asset Management to indicate that an asset can no longer be located. It does not control CMDB CI lifecycle behavior and is not managed by CMDB Data Manager policies.

B. Decommission Incorrect

Decommission describes an operational state of a CI or asset, typically indicating it is no longer in service. However, it is **not a CMDB Data Manager policy type** and does not provide automated end-of-life handling such as retention or restoration.

C. Archive Correct

Archive is an official **CMDB Data Manager end-of-life policy**. It removes CIs from active CMDB use while storing them for a defined retention period. Archived CIs can be **restored if needed**, making this policy useful for compliance and audit requirements.

D. Disposed Incorrect

Disposed is an **asset disposal state** used when hardware or assets are physically destroyed or recycled. It does not apply to CMDB Data Manager policies and does not manage CI lifecycle rules.

E. Retire Correct

Retire is a **CMDB Data Manager end-of-life policy** used when a CI is no longer operational or relevant but must remain in the CMDB for historical, reporting, or audit purposes. Retired CIs are excluded from operational processes while preserving data integrity.

Question: 32

A CMDB Configuration Manager intends to implement CMDB Data Manager delete and archive policies for all

server records in the New York datacenter. In which lifecycle state would servers be affected by these new policies?

- A. End of Life - Retired
- B. In any lifecycle state
- C. Missing - Stolen
- D. Inventory - Available

Answer: A

Explanation:

A. End of Life - Retired

Option-Wise Analysis

A. End of Life - Retired (Correct)

ServiceNow officially positions the **CMDB Data Manager** as a tool that enforces CSDM (Common Service Data Model) standards. According to official documentation and exam guides, **Archive** and **Delete** policies have a built-in safety prerequisite: **CI's must be in the 'Retired' state (or 'End of Life' stage) before they can be archived or deleted.** This prevents the accidental purging of operational or available inventory assets. Even if your policy filters for "New York Datacenter," the processor will only target the CIs within that location that have already transitioned to **Retired**.

B. In any lifecycle state (Incorrect)

While you can set filter conditions like Location or Class, the CMDB Data Manager framework inherently respects lifecycle guardrails. Out-of-the-box, the system prevents "jumping" from an Operational state directly to Deletion. If a server in the New York datacenter is "In Use," the Archive/Delete policies will ignore it until it has been moved through a **Retire** policy first.

C. Missing - Stolen (Incorrect)

"Missing" or "Stolen" are specific sub-states, but they do not satisfy the standard "End of Life" requirement for the automated Archive/Delete cleanup policies in the Data Manager.

D. Inventory - Available (Incorrect)

Records in "Inventory" are considered active assets. Deleting these would cause a major loss of financial and operational visibility. The Data Manager is specifically designed to exclude these from destructive policies.

Question: 33

How does a CMDB Administrator use the ServiceNow Platform to ensure the data quality associated with CIs in the CMDB?

- A. Data Quality Business Rule
- B. Data Quality Scheduled Job
- C. CMDB Workspace
- D. CMDB Audit Business Rule

Answer: D

Explanation:

A. Data Quality Business Rule

– While business rules can enforce specific validations or field-level requirements, they are not designed to monitor or assess overall data quality metrics for CIs. Business rules act immediately on record inserts or updates rather than generating systematic data quality assessments.

B. Data Quality Scheduled Job

– The Data Quality Scheduled Job is the primary mechanism in ServiceNow to automatically evaluate and maintain CMDB data quality. It periodically checks CI records for completeness, correctness, and compliance with organizational standards, generating scores and reports that feed into dashboards like the CMDB Health Dashboard. This allows administrators to proactively monitor and improve data quality across all CIs.

C. CMDB Workspace

– CMDB Workspace is an interface for viewing, managing, and navigating CIs and their relationships. While it helps administrators interact with the CMDB, it does not automatically enforce or measure data quality.

D. CMDB Audit Business Rule

– Audit business rules track changes to CI data for compliance and historical records. They are primarily for recording modifications and supporting audits, not for measuring or improving the overall quality of CI data.

Question: 34

Choose 2 option. The Incident Process Owner asks which classes of CSDM are used on the Incident form. Which classes are appropriate?

- A. Business Application
- B. Application Service
- C. Service Portfolio
- D. Service Offering

Answer: B,D

Explanation:

A. Business Application **Incorrect**

Business Applications represent applications from a business perspective and are mainly used for planning, ownership, and portfolio management. They are **not intended to be selected directly on the Incident form**.

B. Application Service **Correct**

Application Services represent **runtime, operational services** that users and IT teams support. These are appropriate for incidents because they reflect the **actual service being impacted** and allow proper impact analysis and operational reporting.

C. Service Portfolio **Incorrect**

Service Portfolio is a higher-level container used to organize services and offerings. It is **not a CI class** and is not used directly on the Incident form.

D. Service Offering **Correct**

Service Offerings represent **customer-facing or business-facing services** that users interact with. They are appropriate on the Incident form to capture **what service the user is experiencing an issue with**

Question: 35

The ITSM Manager wants to use Technology Management Offerings (Technical Service Offerings) to populate the support group of associated CIs. What CSDM stage would this be completed in?

- A. Walk
- B. Run
- C. Foundation
- D. Crawl
- E. Fly

Answer: A

Explanation:

A. Walk Correct

The **Walk stage** of CSDM adoption focuses on **refining service definitions and linking them to operational details**. Assigning **support groups to CIs** using Technical Service Offerings happens in this stage, enabling ITSM processes such as incident and request management to work effectively. This is the stage where operational support is mapped to services.

B. Run Incorrect

Run is about **operational execution** using the modeled services, like incident, change, and problem management. By the Run stage, support groups are already defined; this stage is about leveraging them operationally.

C. Foundation Incorrect

Foundation is for setting up **core CMDB and organizational data**, such as users, locations, and basic CI classes. It does not include service offerings or mapping support groups.

D. Crawl Incorrect

Crawl establishes **basic service modeling**, like Business Applications and Application Services. It is an initial stage and does not yet include detailed operational associations like support groups.

E. Fly Incorrect

Fly is for **advanced CSDM adoption**, focusing on automation, predictive insights, and optimization. It is beyond the scope of initially populating support groups.

Question: 36

The Application Portfolio Management team is asking for help modeling platforms as Business Applications. How would this be mapped in CSDM using the ServiceNow Platform in the Incident and Change applications?

- A. ServiceNow Platform would be an Architecture type of Platform Host. Incident and Change would be Platform Applications. Then a Depends On relationship would be created between ServiceNow Platform and Incident and Change.
- B. ServiceNow Platform, Incident, and Change would all be web based Architecture type. A Depends On relationship would be created between ServiceNow Platform and Incident and Change.

C. ServiceNow Platform, Incident, and Change would all be web based Architecture type. Incident and Change would have a reference to ServiceNow Platform.

D. ServiceNow Platform would be an Architecture type of Platform Host. Incident and Change would be Platform Applications and have a reference to ServiceNow Platform.

Answer: A

Explanation:

A. ServiceNow Platform would be an Architecture type of Platform Host. Incident and Change would be Platform Applications. Then a Depends On relationship would be created between ServiceNow Platform and Incident and Change Correct

This is the **CSDM best practice**. The ServiceNow Platform is modeled as a **Platform Host** of Architecture type, representing the underlying platform that supports applications. Incident and Change are **Platform Applications**, which run on the platform. A **Depends On relationship** connects these applications to the platform, showing that they rely on it for operation. This modeling ensures accurate operational visibility, supports impact analysis, and aligns with ITSM processes. Simply using a reference field or misclassifying the applications would not meet CSDM standards because it would fail to capture operational dependencies, affecting reporting and service mapping.

B. ServiceNow Platform, Incident, and Change would all be web based Architecture type. A Depends On relationship would be created between ServiceNow Platform and Incident and Change Incorrect

Classifying Incident and Change as Architecture type is incorrect because they are applications, not platforms. This violates CSDM hierarchy and would break proper CMDB modeling.

C. ServiceNow Platform, Incident, and Change would all be web based Architecture type. Incident and Change would have a reference to ServiceNow Platform Incorrect

Using only a reference field does not create a true **Depends On relationship** and does not align with CSDM standards. Operational dependencies would not be correctly captured.

D. ServiceNow Platform would be an Architecture type of Platform Host. Incident and Change would be Platform Applications and have a reference to ServiceNow Platform Incorrect

A reference alone is insufficient. CSDM requires a **Depends On relationship** between Platform Applications and Platform Hosts to ensure proper operational mapping, impact analysis, and ITSM functionality.

Question: 37

A CMDB Administrator is implementing a Vulnerability Response or Security Incident Response and needs to ensure customers have enough context to estimate risk and set task priorities. Which Get Well Playbook from the CSDM Data Foundations Dashboard helps with this?

- A. Application Services with Business Application Relationships
- B. Named Product Models without Product Owners
- C. Locations without a Parent Location
- D. Percentage of Custom Status Values for CI Life Cycle Stages

Answer: A

Explanation:

A. Application Services with Business Application Relationships Correct

This playbook ensures that **Application Services are properly related to Business Applications**, which provides critical business context. That context is essential for **risk estimation, impact analysis, and task prioritization** in Vulnerability Response and Security Incident Response. When security teams understand which business applications are affected, they can prioritize remediation based on business impact.

B. Named Product Models without Product Owners Incorrect

This playbook focuses on asset ownership and accountability. While important for governance, it does not directly provide the **business context needed for risk-based prioritization** in security workflows.

C. Locations without a Parent Location Incorrect

This playbook addresses location hierarchy and data organization. It does not contribute meaningful context for estimating security risk or prioritizing response tasks.

D. Percentage of Custom Status Values for CI Life Cycle Stages Incorrect

This playbook is related to lifecycle data consistency and reporting. It does not help security teams understand **business impact or application context**.

Question: 38

When integrating data into the CMDB using import sets and transform maps, which type of script is added to ensure the data is processed through the IRE?

- A. onComplete
- B. onBefore
- C. onStart
- D. onAfter

Answer: B

Explanation:

B. onBefore

Detailed Reasoning

The use of **onBefore** is the official and technically required method when using the **CMDBTransformUtil** API to bridge Transform Maps with the **IRE (Identification and Reconciliation Engine)**.

How the logic works: The IRE is designed to prevent the standard insert from happening if a duplicate is found. By running the script **onBefore**, the CMDBTransformUtil sends the data to the IRE to find a match. Once the IRE performs the update or creates the new record, the script sets `ignore = true`; This "ignore" command stops the standard transform map from attempting a second, redundant insert.

Why D (onAfter) is technically incorrect for this goal: By the time an **onAfter** script runs, the record has **already been inserted or updated** in the database by the standard transform engine. If you wait until onAfter to call the IRE, you have already bypassed the primary benefit of the IRE: preventing the creation of a duplicate in the first place. You would essentially be updating a record that was just created incorrectly.

Option-Wise Analysis

A. onComplete (Incorrect):

Runs once after the entire data load is finished. It cannot handle row-by-row identification.

B. onBefore (Correct):

The standard event used to intercept data and pass it to the CMDBTransformUtil API, allowing the IRE to manage the record before any database write occurs.

C. onStart (Incorrect):

Runs once before the first row is even touched. It is used for setting variables or clearing logs, not for record identification.

D. onAfter (Incorrect):

While some custom scripts use onAfter for post-processing, it is not the correct place for the IRE de-duplication workflow because the "transformation" (and potential duplicate creation) has already occurred.

Question: 39

A CMDB Administrator has been tasked with gathering information for a presentation to leadership. The Administrator needs to provide Duplicate CI, Orphan CI and Stale CI metrics. Which scorecard provides this information on the CMDB Health Dashboard?

- A. Compliance
- B. Correctness
- C. Completeness

Answer: B

Explanation:**A. Compliance Incorrect**

The **Compliance** scorecard measures adherence to defined standards and policies, such as required attributes and audit rules. It does not focus on duplicates, orphaned, or stale CIs.

B. Correctness Correct

The **Correctness** scorecard evaluates whether CMDB data is **accurate and logically valid**. This includes identifying **duplicate CIs, orphaned CIs, and stale CIs**, making it the appropriate scorecard for these metrics.

C. Completeness Incorrect

The **Completeness** scorecard measures whether required fields and relationships are populated. While important, it does not specifically track duplicates, orphaned, or stale CIs.

Question: 40

A CMDB Administrator is reviewing the health of the CMDB and notices a large percentage of the Hardware CIs are missing serial numbers. The Administrator is concerned this may cause duplicate CIs and would like to resolve the issue in a timely manner. What structured guidelines provided by ServiceNow are available to troubleshoot and resolve the issue?

- A. CMDB Health Dashboard Playbooks
- B. CMDB Data Foundations Dashboard Playbooks
- C. CSDM Data Foundations Dashboard Playbooks

Answer: A

Explanation:

A. CMDB Health Dashboard Playbooks

This is the correct answer. Missing serial numbers on Hardware CIs is a CMDB health issue because serial numbers are key identification attributes used by the Identification and Reconciliation Engine. When they are missing, the risk of duplicate CIs increases and overall CMDB accuracy is reduced. The CMDB Health Dashboard Playbooks provide structured, step-by-step guidance to investigate missing attributes, identify root causes such as discovery or integration gaps, correct data ingestion issues, and prevent future duplicates. These playbooks are specifically designed for troubleshooting and resolving CMDB data quality problems.

B. CMDB Data Foundations Dashboard Playbooks

This option is incorrect because Data Foundations focuses on baseline readiness and coverage of CMDB data, not on actively troubleshooting health issues. While it helps ensure required data exists at a high level, it does not provide prescriptive remediation steps for missing CI attributes or duplicate prevention.

C. CSDM Data Foundations Dashboard Playbooks

This option is incorrect because CSDM Data Foundations playbooks are intended to support CSDM adoption and service modeling readiness. They help ensure that foundational service and application data is in place, but they do not address CI-level attribute quality issues such as missing serial numbers or duplicate hardware records.

D. CSDM Now Create Playbooks

This option is incorrect because Now Create playbooks are part of ServiceNow's implementation methodology. They provide guidance on how to design and implement solutions, but they are not used for operational CMDB health monitoring or data quality remediation.

Question: 41

A CMDB Administrator needs to import external data into the CMDB. As the CMDB Administrator wants to reduce the risk for creating duplicates and to update information from unauthorized sources, it has to be ensured that the Identification and Reconciliation API will not be bypassed. What is the recommended method to import data into the CMDB utilizing the Identification and Reconciliation API?

- A. IntegrationHub ETL
- B. Table API (REST API or SOAP API)
- C. Import Sets and Transform Maps

Answer: A

Explanation:

A. IntegrationHub ETL – IntegrationHub ETL is the recommended method for importing external data into the CMDB without bypassing the Identification and Reconciliation Engine (IRE). It supports deduplication, reconciliation, and automated CI classification, ensuring that imported data adheres to CMDB integrity rules. IntegrationHub ETL is the safest and most robust method for external imports into the CMDB.

B. Table API (REST API or SOAP API) – Using Table APIs to import data can bypass the IRE if not carefully

configured. This increases the risk of creating duplicate CIs or overwriting important data, which is why it is not recommended for CMDB imports when IRE governance is required.

C. Import Sets and Transform Maps – While Import Sets can be used for data import, they do not automatically enforce IRE rules unless combined with additional mapping and reconciliation logic. Out-of-the-box, this method does not guarantee duplicate prevention or full adherence to CMDB governance.

Question: 42

A CMDB Administrator wants only the CIs of Principal Classes to appear in CI reference fields, for example the CI reference fields accessible from an Incident Form. Where does the CMDB Administrator designate Principal Classes?

- A. CMDB Workspace
- B. CI Class Manager
- C. CMDB Data Manager
- D. System Properties

Answer: B

Explanation:

A. CMDB Workspace **Incorrect**

CMDB Workspace provides a centralized interface for viewing, analyzing, and managing configuration items and their relationships. While it helps administrators work with CIs, it does not provide the ability to define or configure which CI classes are marked as Principal Classes.

B. CI Class Manager **Correct**

The **CI Class Manager** is the authoritative place where CI class behavior and metadata are managed. Within the CI Class Manager, on the **Basic Info** tab of a CI class, the CMDB Administrator can designate the class as a **Principal Class**. Principal Classes control which CI classes are exposed and selectable in ITSM processes such as Incident, Change, and Problem, ensuring users interact only with relevant and supported CI types.

C. CMDB Data Manager **Incorrect**

CMDB Data Manager focuses on managing CI lifecycle through policies such as archive, retire, and delete. It does not control CI class visibility or Principal Class designation.

D. System Properties **Incorrect**

System properties control global system behavior and configuration settings. Principal Class configuration is not handled through system properties but through class-level management in the CI Class Manager.

Question: 43

Choose 2 option. A CMDB Administrator wants to run the Services Have Owners Identified playbook to remediate the issues shown in the CMDB Data Foundations Dashboard. Which remediation plays would be used?

- A. Analyze Data
- B. Fix Data
- C. Govern Data

D. Report Data

Answer: B,C

Explanation:

A. Analyze Data is used in some CMDB playbooks to identify data quality issues by examining CI records, relationships, and metrics. However, for the Services Have Owners Identified playbook, the dashboard has already highlighted which services are missing owners. Therefore, additional analysis is not needed in this remediation flow, making Analyze Data irrelevant in this context.

B. Fix Data is the primary remediation play in this playbook. It provides step-by-step guidance to correct missing or inaccurate service owner assignments. Using Fix Data, the administrator can update CI records to ensure that each service has a valid, accurate owner, directly addressing the issues shown on the dashboard.

C. Govern Data ensures that any changes made during remediation comply with organizational policies, approvals, and governance rules. In the context of this playbook, Govern Data guarantees that service owner assignments follow proper authorization procedures and remain consistent with best practices, supporting long-term CMDB data quality and compliance.

D. Report Data is used to generate reports or dashboards for stakeholders, showing CI data health or compliance metrics. While useful for communication, it does not actively remediate or govern data, so it is not part of the remediation process for this playbook.

Question: 44

Choose 2 option. A CMDB Administrator is evaluating whether to monitor the metrics provided on the CMDB Data Foundation Dashboard. Which benefits support the decision to continually monitor the results on this dashboard?

- A. Provides a list of all CIs that failed health audits
- B. Provides metrics on active CIs updated in the last 90 days
- C. Reports on all orphan CIs in the CMDB
- D. Provides metrics for CIs Processed by the IRE

Answer: B,D

Explanation:

A. Provides a list of all CIs that failed health audits Incorrect

This information is provided by the **CMDB Health Dashboard**, not the CMDB Data Foundations Dashboard. Health audits focus on correctness, completeness, and compliance.

B. Provides metrics on active CIs updated in the last 90 days Correct

The CMDB Data Foundations Dashboard includes metrics that show how **recently CI data is being maintained and updated**. Monitoring this helps ensure CI data remains current and relevant, supporting accurate operational and security processes.

C. Reports on all orphan CIs in the CMDB Incorrect

Orphan CI reporting is part of the **CMDB Health Dashboard (Correctness scorecard)**, not the Data Foundations Dashboard.

D. Provides metrics for CIs processed by the IRE Correct

The CMDB Data Foundations Dashboard includes visibility into **CI**s processed through the **Identification and Reconciliation Engine (IRE)**. Continually monitoring this helps ensure proper data ingestion, reduced duplication, and adherence to CMDB governance standards.

Question: 45

A CMDB Administrator would like to minimize stale CIs in the CMDB. Which CMDB Health Dashboard scorecard displays this information?

- A. Compliance
- B. Completeness
- C. Correctness

Answer: C

Explanation:

A. Compliance Incorrect

The Compliance scorecard measures whether CIs meet defined audit rules and policies, such as required attributes and governance standards. It does not track stale data.

B. Completeness Incorrect

The Completeness scorecard focuses on whether required fields and relationships are populated. While important, it does not identify stale or outdated CIs.

C. Correctness Correct

The Correctness scorecard evaluates whether CMDB data is accurate and logically valid. This includes identifying **stale CIs**, **duplicate CIs**, and **orphaned CIs**, making it the appropriate scorecard to monitor and reduce stale records.

Question: 46

A Configuration Manager wants to use the Unified Map. Where would it be accessed?

- A. CMDB Workspace
- B. CI Class Manager
- C. CMDB Data Manager

Answer: A

Explanation:

A. CMDB Workspace Correct

The **Unified Map** is accessed from the **CMDB Workspace**. It provides a consolidated, interactive visualization of CIs and their relationships, allowing Configuration Managers to easily understand how configuration items are connected across the CMDB.

B. CI Class Manager Incorrect

CI Class Manager is used to manage CI class definitions, attributes, and relationships, not to visualize CI

relationships in a map view.

C. CMDB Data Manager Incorrect

CMDB Data Manager focuses on CI lifecycle management through policies such as archive, retire, and delete. It does not provide visualization tools like the Unified Map.

Question: 47

The CMDB Administrator has set-up two Dynamic Reconciliation Rules within the ServiceNow Production Instance. The 'Server' class has a Dynamic Reconciliation Rule of largest value for the RAM field. The 'Windows Server' class

Discovery Source	RAM (MB)
Tivoli	4.096
ServiceNow	4.096
LANDesk	2.048
Altiris	8.192

has a Dynamic Reconciliation Rule of most reported for the RAM field. Given the above data in the Multisource CMDB, which value would be added to the CMDB for RAM for a 'Server' CI?

- A. 8.192 MB
- B. 2.048 MB
- C. 4.096 MB

Answer: A

Explanation:

A. 8.192 MB Correct

Since the CI in question is a **Server** (not specifically a Windows Server), the **Server class dynamic reconciliation rule applies**. The rule for this class is **largest value**, and among the reported values (4.096, 4.096, 2.048, and 8.192 MB), the largest value is **8.192 MB**, which will be written to the CMDB.

B. 2.048 MB Incorrect

This value is neither the largest nor the most frequently reported. It would not be selected by any of the defined dynamic reconciliation rules.

C. 4.096 MB Incorrect

Although this value is the **most reported**, that rule applies only to the **Windows Server** class. Since the CI is evaluated at the **Server** class level, the largest-value rule takes precedence.

Question: 48

Choose 2 option. With CMDB 360/Multisource CMDB the Dynamic Reconciliation Rules will also be enabled. Based on the request of the management, a CMDB Administrator has to set up multiple Dynamic Reconciliation Rules. Which are available 'Dynamic Rule Types' within the 'Create Reconciliation Rule' wizard?

- A. Most Reported
- B. Last Updated
- C. Last Created
- D. Smallest Value

Answer: A,D

Explanation:

A. Most Reported Correct

This rule selects the value that is reported by the greatest number of data sources and is a core Dynamic Reconciliation Rule type.

B. Last Updated Incorrect

“Last Updated” is **not** a Dynamic Reconciliation Rule type. Timestamp-based logic is handled through other mechanisms such as source precedence, not dynamic rules.

C. Last Created Incorrect

“Last Created” is not supported in Dynamic Reconciliation. Reconciliation rules do not evaluate record creation time.

D. Smallest Value Correct

This rule selects the smallest numeric value reported across all sources and is a valid Dynamic Reconciliation Rule type.

Question: 49

A CMDB Administrator identifies duplicate CIs. One was created by a manual import, and the other one was created by automated discovery. The discovered CI has the latest IP address, while the manually imported CI has an accurate relationship to a critical business application. How does the Administrator use the Duplicate CI Remediator to resolve this issue?

- A. Retain the discovered CI, and delete the manually imported CI
- B. Retain the manually imported CI, and delete the discovered CI
- C. Merge the two CIs automatically, retaining all attributes from the discovered CI
- D. Retain the discovered CI, but merge the relationship from the manually imported CI

Answer: D

Explanation:

A. Retain the discovered CI, and delete the manually imported CI Incorrect

Deleting the manually imported CI would remove the **accurate relationships to critical business applications**, which are essential for understanding service dependencies and impact analysis. While the discovered CI has up-to-date attributes, losing these relationships would compromise CMDB completeness and correctness.

B. Retain the manually imported CI, and delete the discovered CI Incorrect

Keeping the manually imported CI would preserve relationships, but the **latest attribute information**, such as the IP address from discovery, would be lost. This reduces the accuracy of the CMDB and can affect automated processes that rely on current CI data.

C. Merge the two CIs automatically, retaining all attributes from the discovered CI Incorrect

Automatically retaining all attributes from the discovered CI would overwrite the manually imported CI's data, including its **relationships to critical business applications**, resulting in **data loss**. This approach ignores the principle of preserving both accurate attributes and relationships.

D. Retain the discovered CI, but merge the relationship from the manually imported CI Correct

The **best practice** is to retain the discovered CI because it has the **most up-to-date attribute information** (IP address, system details) while merging the **critical relationships** from the manually imported CI. The **Duplicate CI Remediator** allows selective merging, so both accurate attributes and important business relationships are preserved. This ensures that the CMDB remains both **accurate and complete**, supporting operational processes, impact analysis, and reporting. Using this approach minimizes duplication while maintaining the integrity and usability of the CMDB.

Question: 50

A CMDB Administrator needs to identify which attributes have been created specifically for the Windows Server class. Which tab in the Attributes section is used?

- A. Child
- B. Added
- C. Derived
- D. All

Answer: B

Explanation:

A. Child Incorrect

The **Child** tab shows attributes inherited from **child classes**, not the attributes specifically created for the class in question.

B. Added Correct

The **Added** tab displays attributes that were **directly created for the selected class** (in this case, Windows Server). This allows administrators to quickly identify class-specific attributes without including inherited or derived fields.

C. Derived Incorrect

The **Derived** tab shows attributes **inherited from parent classes** or other relationships, not the attributes unique to the class itself.

D. All Incorrect

The **All** tab shows every attribute associated with the class, including inherited, derived, and added attributes. While comprehensive, it does not differentiate which attributes were **specifically added** to the Windows Server class.

Question: 51

The CMDB Administrator group aims to display meaningful results on the CMDB Health Dashboard Compliance

Scorecard for server records that are not on the latest patch.What must be configured to achieve this goal?

- A. Technical Service Offerings, Dynamic CI Groups, CMDB Groups
- B. Stale, Orphan, Duplicate
- C. Certification Filter, Certification Template, Audit
- D. Certification Policies, Data Filters, Scheduled.Jobs

Answer: C

Explanation:

A. Technical Service Offerings, Dynamic CI Groups, CMDB Groups – These are used to organize and group CIs for operational management, reporting, or automation. While useful for managing CI collections, they do not define compliance checks or certification rules needed for the Health Dashboard Compliance Scorecard.

B. Stale, Orphan, Duplicate – These are CMDB Health metrics that evaluate data quality issues such as outdated, unlinked, or duplicate CIs. They provide general health scoring but do not measure specific compliance criteria like patch status, so they cannot achieve the goal of showing servers not on the latest patch.

C. Certification Filter, Certification Template, Audit – While these elements relate to CMDB certification, they are individual components and do not fully configure the system to enforce compliance tracking. Certification Filters and Templates define what is measured, and Audits collect data, but without policies and scheduled jobs, the compliance scorecard will not display meaningful, actionable results.

D. Certification Policies, Data Filters, Scheduled Jobs – To achieve the goal, Certification Policies must be defined to specify the criteria (e.g., servers missing the latest patch), Data Filters narrow down which CIs are included in the certification evaluation, and Scheduled Jobs run the certification process regularly to update results. Together, these configurations allow the CMDB Health Dashboard Compliance Scorecard to display meaningful and accurate results for server compliance.

Question: 52

A CMDB Administrator has installed a Service Graph Connector and customized a script transform.What will happen on subsequent upgrades if the default definition of the script transform is updated?

- A. A skipped change is created and no change is made to the script transform definition.
- B. The upgrade stops and reports an error.
- C. The Service Graph Connector upgrade refuses to start.

Answer: A

Explanation:

A. A skipped change is created and no change is made to the script transform definition **Correct**

When a Service Graph Connector is upgraded, **any customized script transforms are preserved**. The upgrade process **creates a skipped change** for the updated default definition but does **not overwrite the administrator's customizations**. This ensures that custom logic continues to function while allowing the upgrade to proceed safely.

B. The upgrade stops and reports an error **Incorrect**

The upgrade does **not fail** or stop due to customized script transforms. ServiceNow's upgrade mechanism is

designed to handle customizations gracefully.

C. The Service Graph Connector upgrade refuses to start Incorrect

The upgrade process will still run; it does not refuse to start. Customizations are simply preserved, and skipped changes are logged for review.

Question: 53

The CMDB Configuration Manager is using the CI Class Manager to manage the group ownership of CI classes and needs to leverage the ownership value specified in the CI Class Manager. When configuring a CMDB Data Manager policy, which group reference field should be set?

- A. Change Group
- B. Approval Group
- C. Support Group
- D. Managed By Group

Answer: D

Explanation:

A. Change Group Incorrect

The **Change Group** is part of ITSM workflows for assigning change tasks. It has no relation to CI class ownership or CMDB Data Manager policies.

B. Approval Group Incorrect

The **Approval Group** controls who approves workflow tasks or policy actions. It does not leverage CI class ownership from CI Class Manager.

C. Support Group Incorrect

The **Support Group** identifies the operational support team for a CI. It is not used for governance or policy enforcement in CMDB Data Manager.

D. Managed By Group Correct

The **Managed By Group** field in CMDB Data Manager policies is explicitly designed to **use the ownership defined in CI Class Manager**. By configuring this field, the policy can automatically target the correct group responsible for the CI class, ensuring governance and accountability during data lifecycle actions such as **retire, archive, or delete**. This is the **standard ServiceNow approach** for aligning CMDB policy enforcement with class-level ownership.

Question: 54

A new custom class is needed to reflect a new application being managed in the CMDB. Which roles are minimally needed to add this custom CI class?

- A. data_classification_admin and personalize_dictionary
- B. sn_cmdb_admin and personalize_dictionary
- C. cmdb_inst_admin and personalize_form
- D. itil_admin and personalize_form

Answer: B

Explanation:

A. data_classification_admin and personalize_dictionary Incorrect

The **data_classification_admin** role is related to data classification and privacy, not CI class creation. **Personalize_dictionary** alone does not grant sufficient privileges to manage CMDB classes.

B. sn_cmdb_admin and personalize_dictionary Correct

The **sn_cmdb_admin** role provides administrative access to the CMDB, including the ability to create or modify CI classes.

Personalize_dictionary allows the user to modify the dictionary entries for the class, which is required when creating a new class with its attributes. Together, these roles are **minimally required** to add a custom CI class.

C. cmdb_inst_admin and personalize_form Incorrect

cmdb_inst_admin manages CI records, not the class structure itself. **Personalize_form** is for customizing forms, which is insufficient to create a new CI class.

D. itil_admin and personalize_form Incorrect

ITIL Admin relates to ITSM process roles, and **personalize_form** only affects form layouts. Neither provides privileges to create or manage CI classes.

Question: 55

Choose 2 option. A CMDB Administrator wants to create a CMDB query to find all databases located in Seattle that are connected to application services. They also want to include incidents related to those databases. Which actions does the company take to build this query?

- A. Add a filter to the database node for location = Seattle
- B. Set the relationship level to 'Up to 2nd level relationships'
- C. Add to the canvas the Incident table from the Non-CMDB Tables list
- D. Add property columns to the application service node

Answer: A,C

Explanation:

A. Add a filter to the database node for location = Seattle Correct

Filtering the **database node** ensures that the query only includes databases physically or logically located in **Seattle**. This is essential to narrow down the dataset to the relevant CIs.

B. Set the relationship level to 'Up to 2nd level relationships' Incorrect

In CMDB Query Builder, you typically **define explicit relationships between nodes** rather than relying on a generic “relationship level.” Setting “Up to 2nd level relationships” is unnecessary if you are connecting the database node directly to the application service node.

C. Add to the canvas the Incident table from the Non-CMDB Tables list Correct

To include incidents related to the databases, the **Incident table** must be added as a **Non-CMDB table** and

linked to the database node. This allows the query to retrieve incident records alongside CI information.

D. Add property columns to the application service node Incorrect

Adding property columns only affects **what fields are displayed** in the query results. It does **not impact the logic** for selecting databases, application services, or incidents.

Question: 56

Which type of CMDB Data Manager policy creates tasks that allow the assigned individual to update fields on the CI record?

- A. Attestation
- B. Compliance
- C. Certification
- D. Audit

Answer: C

Explanation:

- A. Attestation** – Attestation policies are often confused with certification, but in ServiceNow CMDB, attestation verifies ownership or accuracy in a yes/no format. They do not typically create tasks to directly update CI fields.
- B. Compliance** – Compliance policies check whether CIs meet defined business rules or standards. They may generate reports or alerts but do not create tasks for CI updates.
- C. Certification** – Certification policies are designed to review and validate CI data against specific criteria. They can create tasks assigned to individuals, allowing them to update fields on the CI record, such as ownership, lifecycle stage, or other attributes. This is the mechanism that ensures CI data remains accurate and aligned with governance practices.
- D. Audit** – Audit policies track CI changes and monitor adherence but do not generate tasks to update CI records. They are informational rather than actionable for data remediation.

Question: 57

What ensures data volume in the CMDB is manageable?

- A. Scheduled Jobs
- B. Business Rules
- C. Archive Policies

Answer: C

Explanation:

- A. Scheduled Jobs** – Scheduled Jobs automate tasks such as data imports, updates, or reports. While they help maintain CMDB operations, they **do not directly control the total volume of data stored**.
- B. Business Rules** – Business Rules enforce data integrity, validations, and relationships between Configuration Items (CIs). They **ensure correct data entry** but **do not reduce or manage overall data volume**.

C. Archive Policies – Archive Policies are specifically designed to **manage CMDB data volume**. They allow outdated or irrelevant records to be archived or removed according to retention rules, keeping the CMDB **efficient and manageable**.

Question: 58

A CMDB Administrator has built a number of Technology Management Service Offerings (Technical Service Offerings) based on Dynamic CI Groups to better maintain group alignment for the member CI. Which Groups are synced to CIs from the offering that has a relationship to a Dynamic CI Group?

- A. Managed by Group
- B. Approval Group
- C. Owned by Group
- D. Support Group

Answer: A

Explanation:

A. Managed by Group – In ServiceNow CMDB, a **Dynamic CI Group** is a flexible way to group CIs based on certain attributes or criteria. When a **Technical Service Offering** is tied to a Dynamic CI Group, the system can automatically update the “**Managed by**” group on each member CI. This means that whoever is responsible for managing the CI (for example, approving changes, monitoring performance, or maintaining configuration) is **kept in sync with the group defined in the offering**. The advantage is that as CIs are added or removed from the Dynamic CI Group, the Managed by Group on the CI automatically updates, reducing manual work, errors, and misalignments in CI management responsibilities.

B. Approval Group – The Approval Group is used during workflow approvals, for example when a change request needs sign-off. Although it is important for change governance, **it is not linked to Technical Service Offerings for automatic synchronization**. Updating this group is done manually or via separate workflows, not through Dynamic CI Group relationships.

C. Owned by Group – The Owned by Group represents the team or department that owns the CI, often for reporting or financial accountability. This ownership **does not automatically change based on a Technical Service Offering**. Syncing only applies to CI management responsibilities (Managed by Group), not ownership.

D. Support Group – The Support Group indicates the operational support team responsible for incident resolution or operational tasks. While it is a critical field for operations, **Technical Service Offerings do not sync this field from Dynamic CI Groups**, as support responsibilities may differ from management responsibilities.

Question: 59

Which is a purpose or requirement of CMDB Data Manager in ServiceNow?

- A. Automates the archival and deletion of records based on retention policies
- B. Encrypts archived records for enhanced security
- C. Automates the enforcement of relationship rules between CIs in the CMDB

Answer: C

Explanation:

A. Automates the archival and deletion of records based on retention policies – This is the function of **Archive Policies**, not CMDB Data Manager. Archive Policies handle cleanup and lifecycle management to keep the CMDB size manageable.

B. Encrypts archived records for enhanced security – CMDB Data Manager does not handle encryption of archived records. Security and encryption of data are managed through **platform-level encryption and access controls** in ServiceNow.

C. Automates the enforcement of relationship rules between CIs in the CMDB – **CMDB Data Manager** ensures **data integrity and consistency** in the CMDB by enforcing relationship and dependency rules between Configuration Items (CIs). It validates that relationships follow defined rules, prevents incorrect associations, and maintains a reliable and accurate CMDB structure. This is critical for impact analysis, change management, and operational decision-making.

Question: 60

A customer wants recently imported server records to be automatically reclassified into more specific CMDB classes after being discovered using ServiceNow Discovery. During the discovery process, if existing Server [cmdb_ci_server] records are reclassified into the Linux [cmdb_ci_linux_server] and Windows Server [cmdb_ci_win_server] classes, which reclassification operation occurred?

- A. Class Downgrade
- B. Class Switch
- C. Class Upgrade

Answer: C

Explanation:

A. Class Downgrade – A Class Downgrade occurs when a CI is moved from a **more specific class to a more general class**. In this scenario, the Server records are being moved to more specific subclasses, so this is not a downgrade.

B. Class Switch – A Class Switch refers to **changing a CI from one sibling class to another** at the same level, such as switching between two unrelated classes. Here, the reclassification is from a general class to its specific subclasses, so it's not a switch.

C. Class Upgrade – A Class Upgrade occurs when a CI is **moved from a general class to a more specific subclass**. In this case, Server [cmdb_ci_server] records are reclassified into Linux [cmdb_ci_linux_server] and Windows Server [cmdb_ci_win_server] after discovery, which are **more specific classes under the Server parent class**. This ensures that CIs are accurately categorized for management, reporting, and downstream IT processes.

Question: 61

A CMDB Administrator wants to improve data quality related to the CSDM. Which action should the Administrator take to meet this goal?

- A. Start the ServiceNow Health Scan
- B. Use the CSDM Data Foundations Dashboard
- C. Use the default configured CMDB Health Dashboard

Answer: B

Explanation:

A. Start the ServiceNow Health Scan – The ServiceNow Health Scan evaluates **overall instance health**, including customizations, performance, and security, but it does not specifically focus on **CSDM-related data quality** in the CMDB.

B. Use the CSDM Data Foundations Dashboard – The **CSDM Data Foundations Dashboard** is designed specifically to **assess, monitor, and improve CMDB data quality in alignment with the Common Service Data Model (CSDM)**. It provides visibility into CI relationships, classification, and adherence to CSDM principles, helping the administrator identify gaps and improve data accuracy and completeness.

C. Use the default configured CMDB Health Dashboard – The default CMDB Health Dashboard measures **general CI health, completeness, correctness, and duplicates**, but it is **not tailored for CSDM-specific alignment**. While useful for overall CMDB health, it does not provide the targeted insights needed for CSDM data quality improvements.

Question: 62

A CMDB Manager wants to improve data quality using the CMDB Health Dashboard. What needs to happen to generate CMDB health scores?

- A. The scheduled jobs for the CMDB Health Dashboard must be activated
- B. The plugin, CMDB health calculation, needs to be installed
- C. Nothing, CMDB health scores are calculated by default

Answer: A

Explanation:

A. The scheduled jobs for the CMDB Health Dashboard must be activated – CMDB health scores are **not calculated automatically by default**. To generate scores for **completeness, correctness, and compliance**, the CMDB Manager must ensure that the **scheduled jobs for the CMDB Health Dashboard** are activated. These jobs run periodically to evaluate CI data and calculate health metrics, which are then displayed on the dashboard for monitoring and improvement purposes.

B. The plugin, CMDB health calculation, needs to be installed – No separate plugin installation is required for CMDB health scoring. The CMDB Health Dashboard functionality comes **pre-installed with the platform**; only the scheduled jobs need to be enabled.

C. Nothing, CMDB health scores are calculated by default – CMDB health scores are **not automatically calculated**. Without activating the scheduled jobs, the dashboard will not display any metrics, and data quality improvements cannot be monitored effectively.

Question: 63

Choose 2 options. Configuration Management requires an accurate inventory of devices to be reflected in the CMDB. Which are common use cases for using Agent Client Collector (ACC)?

- A. Network devices in the DMZ
- B. Devices in secure environments

- C. Devices that intermittently connect to the network
- D. Servers in the data center

Answer: B,C

Explanation:

A. Network devices in the DMZ – Network devices in the DMZ (Demilitarized Zone) are typically discovered using **network-based discovery protocols** like SNMP, not ACC. ACC is agent-based and is better suited for endpoints that require direct client access.

B. Devices in secure environments – ACC is ideal for devices located in **restricted or secure environments** where direct network discovery might be blocked by firewalls or security policies. The agent installed on these devices can collect inventory data and report it back to the CMDB.

C. Devices that intermittently connect to the network – Devices that do not stay continuously connected (like laptops or mobile devices) are best handled using ACC. The agent can gather data **even during brief connections**, ensuring the CMDB has up-to-date information.

D. Servers in the data center – Servers in a controlled data center environment are usually **always on and connected**, making them suitable for **network-based discovery or scheduled scans** rather than ACC. ACC is not necessary for these devices.

Question: 64

A CMDB Data Manager needs to access the ServiceNow platform to create, publish, and manage policies that automate and govern CI lifecycle operations, ensuring the CMDB remains healthy and efficient. Where can the Data Manager do this?

- A. CMDB Workspace Management tab
- B. Service Operations Workspace
- C. CMDB Workspace CMDB 360 tab
- D. CI Class Manager

Answer: A

Explanation:

A. CMDB Workspace Management tab – The **Management tab within CMDB Workspace** is the primary location where a CMDB Data Manager can **create, publish, and manage CMDB governance policies**. These policies automate CI lifecycle operations such as classification, reconciliation, archival, and data quality enforcement. This centralized workspace is specifically designed to keep the CMDB **healthy, efficient, and aligned with best practices**.

B. Service Operations Workspace – Service Operations Workspace focuses on **operational monitoring, incidents, problems, and service health**, not on defining or governing CMDB lifecycle policies. It does not provide tools for CMDB policy creation or management.

C. CMDB Workspace CMDB 360 tab – The **CMDB 360 tab** provides a **holistic, read-only view of CI relationships, dependencies, and insights**. While it is useful for analysis and visualization, it is **not used to create or manage CMDB governance policies**.

D. CI Class Manager – CI Class Manager is used to **define and maintain the CMDB class hierarchy and attributes**. It does not handle lifecycle automation, governance policies, or ongoing CMDB health management.

Question: 65

Choose 2 option. A CMDB Administrator is considering whether to start using the playbooks provided on the CMDB Data Foundation Dashboard. What are the benefits to support the decision to leverage this feature?

- A. Offers insight into the downstream impacts of poorly performing metrics
- B. Offers remediation options to address and improve poorly performing metrics
- C. Offers automated scripts to resolve poorly performing metrics
- D. Offers remediation templates to improve poorly performing metrics

Answer: A,B

Explanation:

A. Offers insight into the downstream impacts of poorly performing metrics

B. Offers remediation options to address and improve poorly performing metrics

Final Option-Wise Analysis

A. Offers insight into downstream impacts (Correct):

The playbooks bridge the gap between "technical data" and "business value." For every metric (like "CIs with no assigned company"), the playbook provides a clear explanation of how this failure affects **downstream** processes like Integrated Risk Management (IRM), Incident routing, and financial reporting. This provides the "why" behind the remediation.

B. Offers remediation options (Correct):

This is a core feature of the playbooks. They provide the "how-to" by outlining specific **remediation options**. While they are not "one-click" buttons, they provide the guided steps and recommended approaches (e.g., "Adjust your Identification Rules" or "Update the Discovery Schedule") required to improve the score.

C. Offers automated scripts (Incorrect):

ServiceNow prioritizes safety in the CMDB. The playbooks do not contain "one-click" scripts to fix data because automated mass-updates could cause significant damage if the root cause isn't understood. All remediation is intended to be a deliberate, governed process.

D. Offers remediation templates (Incorrect):

In ServiceNow terminology, "Templates" refers to reusable configuration objects (like a Propose Change template or a Remediation Rule template in CMDB Health). The Data Foundation playbooks are **instructional guides**, not configuration templates.

Question: 66

A CMDB Administrator, viewing the CMDB Data Foundations Dashboard, notices the Unique Locations Result percentage low. What is the recommended process from the associated playbook to correct this issue?

- A. Use the Duplicate CI Remediator to merge the duplicate location records
- B. Keep both locations as either can be used as a valid alternate location
- C. Review both locations, update CIs with the correct location and delete the duplicate location

D. Retain the location that matches the organization's standard naming convention, and delete the duplicate without further validation

Answer: C

Explanation:

A. Use the Duplicate CI Remediator to merge the duplicate location records – The Duplicate CI Remediator is designed for **Configuration Items**, not for **reference data like Location records**. Location duplicates require validation and cleanup rather than automated CI merging.

B. Keep both locations as either can be used as a valid alternate location – Keeping duplicate locations continues to **violate the uniqueness requirement** and does not improve the Unique Locations metric. This approach leaves CMDB data quality issues unresolved.

C. Review both locations, update CIs with the correct location and delete the duplicate location – The playbook recommends **validating which location is correct**, updating all affected CIs to reference that correct location, and then **removing the duplicate record**. This ensures location data remains accurate, unique, and aligned with CMDB best practices.

D. Retain the location that matches the organization's standard naming convention, and delete the duplicate without further validation – Naming standards alone are **not sufficient validation**. The playbook emphasizes reviewing and confirming accuracy before deleting records to avoid data loss or incorrect CI references.

Question: 67

A CMDB Administrator needs insights into how their CMDB is configured according to ServiceNow recommended practice. Which should be used?

- A. CMDB Data Foundation Dashboard
- B. CMDB Workspace
- C. CMDB Data Manager
- D. CMDB Health Dashboard

Answer: A

Explanation:

A. CMDB Data Foundation Dashboard

The **CMDB Data Foundation Dashboard** is specifically designed to provide insight into how the CMDB aligns with **ServiceNow recommended practices**, including **foundational data standards, CSDM alignment, and configuration expectations**. It evaluates whether key data elements and relationships are structured according to ServiceNow guidance, making it the correct tool for understanding **recommended-practice configuration**.

B. CMDB Workspace

CMDB Workspace is an **operational interface** for managing CIs, relationships, and policies. It does not assess or report on alignment with ServiceNow recommended practices.

C. CMDB Data Manager

CMDB Data Manager is used to **define and enforce governance policies** (classification, lifecycle, reconciliation), not to provide insight or assessment of whether the CMDB is configured according to best

practices.

D. CMDB Health Dashboard

The CMDB Health Dashboard focuses on **data quality metrics** such as completeness, correctness, and compliance. It evaluates **data health**, not whether the CMDB configuration follows ServiceNow recommended practices.

Question: 68

A Configuration Management team has decided to start taking advantage of the CMDB 360/Multisource CMDB functionality. Which system property must be enabled?

- A. glide.identification_engine.multisource.query.max.limit
- B. glide.identification_engine.multisource_non_cmdb_ci_enabled
- C. glide.identification_engine.multisource_cmdb_ci_enabled
- D. glide.identification_engine.multisource_enabled

Answer: D

Explanation:

A. glide.identification_engine.multisource.query.max.limit

This property controls the **maximum number of results returned** during multisource queries. It is a tuning parameter and **does not enable Multisource CMDB functionality**.

B. glide.identification_engine.multisource_non_cmdb_ci_enabled

This property applies only to **non-CMDB CI records** once multisource is already enabled. It cannot activate multisource functionality on its own.

C. glide.identification_engine.multisource_cmdb_ci_enabled

This property controls multisource behavior specifically for **CMDB CI tables**, but it still **depends on the core multisource capability being enabled first**.

D. glide.identification_engine.multisource_enabled

This is the **primary system property** that enables **Multisource CMDB and CMDB 360 functionality**. Once enabled, the Identification and Reconciliation Engine (IRE) can manage **multiple data sources per CI**, allowing CMDB 360 views and improved data governance.

Question: 69

Where can an administrator perform Natural Language Queries (NLQ)?

- A. CMDB Health Dashboard
- B. CMDB Workspace
- C. CMDB Data Manager
- D. CI Class Manager

Answer: B

Explanation:

- A. CMDB Health Dashboard** – The CMDB Health Dashboard is used to **monitor data quality metrics** such as completeness, correctness, and compliance. It does not support searching or querying data using natural language.
- B. CMDB Workspace** – **CMDB Workspace** supports **Natural Language Queries (NLQ)**, allowing administrators to search the CMDB using plain-language questions like “Show all Linux servers running in production”. This feature improves usability and makes it easier to explore CIs, relationships, and dependencies without building complex queries.
- C. CMDB Data Manager** – CMDB Data Manager focuses on **policy creation and governance** for CI lifecycle management. It does not provide search or query capabilities using natural language.
- D. CI Class Manager** – CI Class Manager is used to **define and maintain the CMDB class hierarchy and attributes**, not to search or query CMDB data.

Question: 70

A CMDB Administrator aims to utilize CSDM life cycle field mappings to better align with CSDM best practices. What is the next step to take after selecting the Enable Life Cycle Sync button?

- A. Fix the incorrect values in the Life Cycle Stage to match legacy values
- B. Resolve any incomplete field mappings identified in the Discrepancy Report
- C. Activate the CSDM Life Cycle field mappings

Answer: B

Explanation:

A. Fix the incorrect values in the Life Cycle Stage to match legacy values

CSDM lifecycle synchronization is designed to **standardize data to CSDM**, not adjust CSDM fields to match legacy values. Legacy cleanup is handled later, not as the immediate next step.

B. Resolve any incomplete field mappings identified in the Discrepancy Report

After clicking **Enable Life Cycle Sync**, ServiceNow generates a **Discrepancy Report**.

This report highlights **missing, incomplete, or conflicting field mappings** between legacy lifecycle fields and CSDM lifecycle fields.

The **required next step** is to **review and resolve these incomplete mappings** before lifecycle synchronization can be safely activated.

C. Activate the CSDM Life Cycle field mappings

Activation **cannot occur until all discrepancies are resolved**. Activation is the **final step**, not the immediate next step after enabling lifecycle sync.

Question: 71

A health organization must track certain data (for example, regulated patient information) and its relation to Business Applications. Which action does CSDM recommend to meet this goal?

- A. Create an Information Object to represent the patient information, and then link it through a relationship to the Business Application after consulting with the Application owner.
- B. Create fields on the Business Application record to mark the Business Application as containing patient information, and then ask the Business Application owner to mark the application as having patient information or not.
- C. Work with the Database administration team to classify the data on each database that holds patient information, and then use Relationships to map that back to the Business Application.

Answer: A

Explanation:

A. ✓ Create an Information Object to represent the patient information, and then link it through a relationship to the Business Application after consulting with the Application owner.

This aligns with CSDM best practices. CSDM recommends modeling regulated or critical data as **Information Objects**, which represent logical groupings of information such as patient data. These Information Objects are then related to Business Applications using defined relationships, preserving clear ownership, traceability, and governance without overloading application records with data-specific attributes.

B. Create fields on the Business Application record to mark the Business Application as containing patient information, and then ask the Business Application owner to mark the application as having patient information or not.

This is not recommended by CSDM. Adding custom flags or fields to Business Application records mixes data classification concerns into the application model, reduces consistency, and limits reuse and scalability across the CMDB.

C. Work with the Database administration team to classify the data on each database that holds patient information, and then use Relationships to map that back to the Business Application.

This approach is too implementation-focused and operates at the infrastructure layer. CSDM emphasizes modeling information at a conceptual level (Information Objects) rather than relying on physical database classifications, which may change independently of business applications.

Question: 72

A Configuration Management Governance team is transitioning from utilizing legacy CMDB status fields to CSDM life cycle status fields. Which table can be modified?

- A. Life Cycle Controls [life_cycle_control]
- B. Life Cycle Stage Status [life_cycle_stage_status]
- C. Life Cycle Stages [life_cycle_stage]
- D. Life Cycle Mapping [life_cycle_mapping]

Answer: D

Explanation:

A. Life Cycle Controls [life_cycle_control]

This table is system-controlled and defines governance rules around lifecycle behavior. It is not intended to be modified as part of a transition from legacy CMDB status fields.

B. ✓ Life Cycle Stage Status [life_cycle_stage_status]

This is correct. CSDM allows modification of the Life Cycle Stage Status table so organizations can align lifecycle statuses (such as Installed, In Service, Retired) with their governance and operational needs when moving away from legacy CMDB status fields.

C. Life Cycle Stages [life_cycle_stage]

This table defines the high-level lifecycle stages (for example, Plan, Build, Operate, Retire). These stages are standardized in CSDM and should not be modified.

D. Life Cycle Mapping [life_cycle_mapping]

This table is used to map legacy status values to lifecycle statuses during migration. It supports the transition but is not the table where lifecycle definitions themselves should be modified.

Question: 73

Choose 2 option.What types of policies can be created within CMDB Data Manager?

- A. Retire
- B. De-duplication
- C. Archive
- D. Reconciliation

Answer: A,C

Explanation:

A. Retire

C. Archive

Final Option-Wise Analysis

A. Retire (Correct):

This is one of the four core policy types in the **CMDB Data Manager** (the others being Archive, Delete, and Attest). A Retire policy allows you to automate the transition of CIs to a "Retired" lifecycle state based on specific conditions, such as inactivity or an upcoming decommissioning date.

B. De-duplication (Incorrect):

While de-duplication is vital for CMDB health, it is **not** a policy type within the CMDB Data Manager. De-duplication is managed by the **IRE (Identification and Reconciliation Engine)** via **Identification Rules** and resolved using the **Duplicate CI Remediator** wizard.

C. Archive (Correct):

The CMDB Data Manager specifically includes **Archive** as a policy type. This allows administrators to move records that have already reached the "End of Life" stage into archive tables. This keeps the active CMDB high-performing while retaining data for historical or compliance reasons.

D. Reconciliation (Incorrect):

Reconciliation is a core CMDB function, but it is not a "policy" inside the Data Manager. It is configured via **Reconciliation Rules** in the **CI Class Manager**, where you define which discovery sources (like Qualys vs. ServiceNow Discovery) have priority for specific CI attributes.

Question: 74

An organization utilizes multiple data sources to update its CMDB, each assigned a different priority level. A high-priority data source is scheduled to update server records weekly. However, due to an integration issue, the high-priority data source stops updating the records. Which configuration can be used to allow a lower-priority data source to update records after a specified period of inactivity from the higher-priority source?

- A. Reconciliation Rules
- B. Data Refresh Rules
- C. Identification Rules
- D. Health Inclusion Rules

Answer: B

Explanation:

A. Reconciliation Rules

Reconciliation rules control **which data source wins when multiple sources update the same attribute at the same time**. They are based on source priority, not time. If the higher-priority source stops sending updates, reconciliation rules alone will still prevent lower-priority sources from updating those attributes.

B. ✓ Data Refresh Rules

This is correct. Data Refresh Rules define **how long data from a source remains authoritative**. If a high-priority data source has not updated a CI or attribute within a configured time window, the data is considered stale. At that point, **lower-priority data sources are allowed to update the record**, ensuring the CMDB does not become outdated due to inactive integrations.

C. Identification Rules

Identification rules determine **how CIs are identified and matched** (for example, which attributes uniquely identify a server). They do not control update precedence or timing between data sources.

D. Health Inclusion Rules

Health Inclusion Rules define **which CIs are included in CMDB Health scoring**. They have no impact on data source priority, update behavior, or fallback logic.

Question: 75

The CMDB Configuration Management team wants to manage de-duplication tasks generated from data ingested into the CMDB via the Identification and Reconciliation Engine (IRE). In which area of the CMDB Workspace can they locate these de-duplication tasks?

- A. Total status tile under the My Work tab
- B. CMDB feature adoption tile under the Insights tab
- C. Important actions tile under the Home tab

Answer: C

Explanation:

C. Important actions tile under the Home tab

Detailed Reasoning

The **CMDB Workspace** is designed to centralize high-priority configuration management tasks.

The Home Tab Logic: The **Home** view serves as the primary "landing page" for the CMDB Administrator. It is specifically built to surface "at-a-glance" health metrics and immediate work items.

Important Actions Tile: This specific tile (often appearing as the "Health-related tasks" or "Important Actions" card) is where the system flags **Duplicate CIs**, Orphan CIs, and Stale CIs. It acts as a prioritized "to-do" list. Clicking the **Duplicate CIs** card within this tile provides the most direct route to the list of de-duplication tasks generated by the IRE.

Option-Wise Analysis

A. Total status tile under the My Work tab (Incorrect)

The "My Work" tab is primarily for managing tasks assigned specifically to you or your group via **CMDB Data Manager** workflows (like attestation or retirement tasks). While you can find work there, the **Home** tab's "Important Actions" is the designated entry point for health-driven de-duplication tasks.

B. CMDB feature adoption tile under the Insights tab (Incorrect)

The **Insights** tab is a reporting and analytics area. It shows you how well you are using the CMDB (e.g., your CSDM maturity or your coverage of the IRE). It is not an operational area for managing daily tasks.

C. Important actions tile under the Home tab (Correct)

This is the official "Command Center" for the Configuration Manager. It summarizes the most critical issues discovered by the CMDB Health Dashboard and the IRE, providing direct links to remediate duplicates.

Question: 76

A CMDB Administrator is creating technical documentation for stakeholders, which includes a list of attributes, Identification and Reconciliation Engine (IRE) rules, and suggested relationships for several classes. Which central location does the CMDB Administrator use to collect this information?

- A. CMDB Data Manager
- B. CI Class Manager
- C. CI Identifiers
- D. CMDB Workspace

Answer: B

Explanation:

A. CMDB Data Manager

CMDB Data Manager focuses on data governance activities such as reconciliation, de-duplication, data refresh, and source prioritization. It does not serve as a centralized reference for class-level attributes, IRE rules, and recommended relationships.

B. ✓ CI Class Manager

This is correct. CI Class Manager is the central place where CMDB administrators can view and manage **CI class definitions**, including attributes, Identification and Reconciliation Engine (IRE) identification rules, reconciliation behavior, and suggested or recommended relationships. It is commonly used as the

authoritative reference when documenting how classes are modeled and governed.

C. CI Identifiers

CI Identifiers is focused specifically on the configuration of identification rules used by the IRE. It does not provide a complete view of attributes or recommended relationships for multiple classes.

D. CMDB Workspace

CMDB Workspace is designed for operational oversight, health, and remediation activities. While it provides insights and tasks, it is not the primary location for reviewing or documenting class definitions and modeling details.

Question: 77

A CMDB Manager uses CMDB 360/Multisource CMDB to maintain and improve CMDB quality. Why would the Manager use CMDB 360/Multisource CMDB?

- A. To identify CI attributes from multiple data sources
- B. To ingest data from multiple data sources using Import Set(s)
- C. To ingest data from multiple data sources using Service Graph Connector(s)
- D. To populate the CMDB from multiple data sources

Answer: A

Explanation:

A. ✓ To identify CI attributes from multiple data sources

This is correct. CMDB 360 (also known as Multisource CMDB) is used to **analyze, compare, and visualize CI data coming from multiple sources**. It allows CMDB managers to see which attributes are populated by which data sources, assess data completeness and trust, and identify gaps or conflicts. This supports ongoing CMDB quality improvement and governance.

B. To ingest data from multiple data sources using Import Set(s)

This is incorrect. Import Sets are a data ingestion mechanism, but CMDB 360 does not ingest data. It operates on data that has already been loaded into the CMDB.

C. To ingest data from multiple data sources using Service Graph Connector(s)

This is incorrect. Service Graph Connectors are used for ingesting and normalizing data from external systems into the CMDB. CMDB 360 does not perform ingestion; it analyzes and visualizes ingested data.

D. To populate the CMDB from multiple data sources

This is incorrect. Populating the CMDB is done by discovery tools, integrations, Import Sets, or Service Graph Connectors. CMDB 360 is used **after ingestion** to evaluate and improve the quality of the data.

Question: 78

A CMDB Administrator needs to ingest relevant data from Microsoft SCCM into the CMDB. Which ingestion method brings the fastest time to value?

- A. Agent Client Collector
- B. Service Graph Connectors
- C. IntegrationHub ETL
- D. Import Sets

Answer: B

Explanation:

A. Agent Client Collector

This is incorrect. Agent Client Collector requires deploying agents and configuring custom data collection. It takes more setup time and is not optimized for rapid ingestion from enterprise systems like SCCM.

B. ✓ Service Graph Connectors

This is correct. Service Graph Connectors provide **out-of-the-box, prebuilt integrations** for common data sources such as Microsoft SCCM. They include predefined data mappings, normalization, and Identification and Reconciliation Engine (IRE) support, delivering the **fastest time to value** with minimal configuration.

C. IntegrationHub ETL

This is incorrect. While powerful and flexible, IntegrationHub ETL requires designing custom flows, transformations, and mappings. This increases implementation time compared to ready-made Service Graph Connectors.

D. Import Sets

This is incorrect. Import Sets are a low-level ingestion mechanism that typically require custom scripts, transform maps, and reconciliation logic. They offer flexibility but not fast time to value.

Question: 79

A CMDB Administrator group aims to establish a process for receiving task notifications when the Support Group or Managed By Group fields are not populated for operational Linux servers stored in the CMDB. Which ServiceNow modules can be leveraged to configure recommended fields and generate task records in cases where these fields are missing for Linux servers in the CMDB?

- A. Technical Service Offerings and Dynamic CI groups
- B. CMDB Workspace and Scheduled Jobs
- C. CI Class Manager and Health Preferences
- D. Dynamic CI groups and CMDB groups

Answer: C

Explanation:

A. Technical Service Offerings and Dynamic CI groups

Technical Service Offerings relate to service modeling, not CMDB data quality enforcement. Dynamic CI groups can group CIs but do not define required/recommended fields or generate remediation tasks when attributes are missing.

B. CMDB Workspace and Scheduled Jobs

CMDB Workspace is used to view and manage tasks and health insights, but it does not define recommended

fields. Scheduled Jobs are not used to enforce CMDB field population or generate health-based tasks.

C. ✓ CI Class Manager and Health Preferences

This is correct.

CI Class Manager is used to define **recommended (and mandatory) attributes** for specific CI classes, such as Linux servers.

Health Preferences are used to configure **CMDB Health rules**, including completeness checks for missing attributes like Support Group or Managed By. When these fields are missing, CMDB Health can automatically **generate task records** for remediation and notify the appropriate teams.

D. Dynamic CI groups and CMDB groups

These are used for organizing and filtering CIs, not for defining attribute requirements or triggering CMDB Health tasks.

Question: 80

A data center has many servers. The CMDB Administrator wants to confirm that all servers exist. Which Data Manager policy type does the Administrator implement?

- A. Promotion
- B. Verification
- C. Attestation
- D. Certification

Answer: B

Explanation:

A. Promotion

Promotion policies are not used to confirm the existence of Configuration Items. Promotion is related to moving data between environments or states, not validating whether servers actually exist.

B. ✓ Verification

This is correct. **Verification policies** in CMDB Data Manager are used to confirm that CIs (such as servers) actually exist based on authoritative data sources. In this scenario, the administrator wants to verify that all servers in the data center are present and accounted for in the CMDB, which is exactly the purpose of verification.

C. Attestation

Attestation involves asking owners or stakeholders to manually confirm the accuracy of CI data (for example, via periodic reviews). It is a human-driven validation process, not an automated existence check.

D. Certification

Certification is a formal governance process used to validate CMDB data against defined standards, often for compliance or audit purposes. It is broader than simply confirming that servers exist.

Question: 81

The Configuration Management team wants to confirm that all servers in the CMDB actually exist in the data center. Which CMDB Data Manager policy type would the team create?

- A. Retire
- B. Attestation
- C. Delete
- D. Certification
- E. Archive

Answer: B

Explanation:

A. Retire

Retire policies are used to transition Configuration Items out of service when they are no longer needed. They do not validate whether servers currently exist in the data center.

B. ✓ Attestation

This is correct. **Attestation policies** are used to ask a **human owner or responsible group** to confirm that a CI is valid and actually exists. In this scenario, the team wants to confirm that all servers in the CMDB physically exist in the data center, which is exactly the purpose of attestation.

C. Delete

Delete actions remove records from the CMDB. They are not used to confirm existence and are typically the result of governance decisions, not a validation process.

D. Certification

Certification is a broader governance and compliance process that validates CMDB data against defined standards. It is more comprehensive than simply confirming the existence of servers.

E. Archive

Archive policies are used for data retention and historical purposes. They do not validate whether CIs currently exist.

Question: 82

A CMDB Administrator has a number of similar de-duplication tasks that need to be remediated in bulk. How does the Administrator achieve this?

- A. Utilize the Duplicate CI Remediator Wizard
- B. Create and run a de-duplication template
- C. Configure and run a custom de-duplication background script
- D. Create de-duplication tasks manually and remediate each

Answer: B

Explanation:

A. Utilize the Duplicate CI Remediator Wizard

The Duplicate CI Remediator Wizard is designed for **interactive, one-off or small numbers of duplicate CIs**. It is not efficient for handling a large number of similar de-duplication tasks in bulk.

B. ✓ Create and run a de-duplication template

This is correct. **De-duplication templates** allow CMDB Administrators to define reusable remediation logic (such as which CI to retain, merge rules, and attribute handling) and then apply it to **many similar de-duplication tasks at once**, enabling efficient bulk remediation.

C. Configure and run a custom de-duplication background script

Custom scripts are not a recommended or supported approach for CMDB de-duplication. They bypass CMDB governance and can introduce data integrity risks.

D. Create de-duplication tasks manually and remediate each

This is inefficient and defeats the purpose of bulk remediation. CMDB provides built-in mechanisms specifically to avoid manual, one-by-one resolution.

Question: 83

A Windows server is reclassified from the Server table [cmdb_ci_server] to the Windows Server table [cmdb_ci_win_server] when processed through the Identification and Reconciliation Engine (IRE). Which process occurred?

- A. Class Downgrade
- B. Class Upgrade
- C. Class Change
- D. Class Switch

Answer: B

Explanation:

B. Class Upgrade

Final Option-Wise Analysis

In the ServiceNow CMDB and the **Identification and Reconciliation Engine (IRE)**, the movement between classes is strictly categorized based on the direction within the class hierarchy:

B. Class Upgrade (Correct):

This is the official term for moving a CI **from a parent class to a more specific child class**. When your Qualys scan identifies that a generic cmdb_ci_server is actually running a Windows OS, the IRE "Upgrades" it to the cmdb_ci_win_server table. This is the most common path during discovery as more granular data is collected.

A. Class Downgrade (Incorrect):

A downgrade is the exact opposite. It occurs when a CI moves **from a specific child class to a more generic parent class**. This might happen if a discovery source that previously provided OS details is replaced by a source that only knows the device is a server.

D. Class Switch (Incorrect):

In ServiceNow's technical framework, a **Class Switch** is the term used specifically for moving a CI **between sibling classes** (classes that share the same parent but are not parents/children of each other). For example,

moving a CI from the Linux Server table to the Windows Server table is a Class Switch.

Question: 84

Choose 2 option. A Windows administration team wants a grouping of CIs using CMDB groups. Which methods can be used?

- A. Encoded queries
- B. Saved queries
- C. Scripted queries
- D. Tag-based queries

Answer: A,B

Explanation:

A. Encoded queries

B. Saved queries

Detailed Option Analysis

ServiceNow defines three specific "Group Content" types that can be added to a CMDB Group. Options A and B are two of those core types.

A. Encoded queries (Correct):

This is the most common method. You use the standard condition builder (e.g., Class is Windows Server AND Location is New York) to dynamically include CIs. As soon as a new Windows server is discovered in that location, it is automatically added to the group.

B. Saved queries (Correct):

CMDB Groups allow you to leverage **CMDB Query Builder** results. This is the only way to group CIs based on **relationships**. For example, if the Windows team wants a group of "All Windows Servers that host a Production SQL Database," you would create that relationship query in the Query Builder, save it, and then add that **Saved Query** to the CMDB Group.

C. Scripted queries (Incorrect):

While ServiceNow supports scripts in many areas, "Scripted Query" is not an out-of-the-box population method for a CMDB Group. Advanced logic is instead handled via the **CMDB Query Builder** (Saved Queries).

D. Tag-based queries (Incorrect):

"Tag-based queries" is not a recognized population method for CMDB Groups. If you wanted to group by tags, you would use an **Encoded Query** (Option A) to filter the `discovery_tag` table or a specific tag field.

Question: 85

A CMDB Administrator knows that the CMDB Data Foundation Dashboard is a resource to monitor and improve data quality. What is a benefit of this dashboard?

- A. Provides the ability to configure health-related metrics

- B. Provides key health-related metrics to make decisions
- C. Provides the ability to resolve certification policy tasks

Answer: B

Explanation:

A. Provides the ability to configure health-related metrics

This is incorrect. The CMDB Data Foundation Dashboard is a **reporting and visibility tool**, not a configuration interface. Health rules, metrics, and preferences are configured elsewhere (for example, in Health Preferences and CI Class Manager).

B. ✓ Provides key health-related metrics to make decisions

This is correct. The CMDB Data Foundation Dashboard gives stakeholders **visibility into CMDB data quality**, such as completeness, correctness, compliance, and overall health trends. These metrics help CMDB administrators and governance teams identify problem areas, prioritize remediation, and track improvements over time.

C. Provides the ability to resolve certification policy tasks

This is incorrect. Certification tasks are resolved through CMDB Data Manager workflows and task records, not through the Data Foundation Dashboard. The dashboard may highlight issues, but it does not provide task resolution capabilities.

Question: 86

The CMDB Administrator wants to leverage the Staleness metric from the CMDB Health Dashboard - Correctness Scorecard. Which field is used to calculate the duration of this metric?

- A. First discovered (first_discovered)
- B. Created (sys_created_on)
- C. Most recent discovery (last_discovery)
- D. Last modified on (last_modified)
- E. Updated (sys_updated_on)

Answer: C

Explanation:

A. First discovered (first_discovered)

This field records when the CI was initially discovered for the first time. It does not change over time and therefore cannot be used to measure how stale the data has become.

B. Created (sys_created_on)

This field indicates when the CI record was created in the CMDB. It reflects record creation, not ongoing data freshness or discovery activity.

C. ✓ Most recent discovery (last_discovery)

This is correct. The **Staleness metric** in the CMDB Health Dashboard's Correctness Scorecard is calculated based on the **last time the CI was successfully discovered**. The longer the duration since the last_discovery timestamp, the more stale the CI is considered.

D. **Last modified on (last_modified)**

This field captures the last time any update occurred on the record, including manual changes. It does not reliably indicate whether the CI data is still accurate from a discovery perspective.

E. **Updated (sys_updated_on)**

Similar to last_modified, this field reflects any update to the record, not specifically discovery activity, and therefore cannot be used to calculate staleness.

Question: 87

A CMDB Administrator needs to track which CIs and CI classes are missing key data. Which CMDB Health Dashboard scorecard supports tracking this requirement?

- A. Correctness
- B. Compliance
- C. Completeness

Answer: C

Explanation:

A. **Correctness**

Correctness focuses on whether CI data is **accurate and up to date**, such as staleness based on last discovery. It does not track whether required or recommended fields are missing.

B. **Compliance**

Compliance measures whether CIs conform to defined policies, standards, or governance rules. While important, it is not specifically focused on missing attribute values.

C. ✓ **Completeness**

This is correct. The **Completeness scorecard** tracks whether **required or recommended attributes are populated** for CIs and CI classes. It helps identify which CIs and classes are missing key data fields and where remediation is needed.

Question: 88

A CMDB Administrator is asked to clean up the CMDB duplicates. What is the preferred way to manage this task?

- A. The de-duplication task module
- B. The de-duplication dashboard on the CMDB workspace
- C. My Tasks in the Application Navigator

Answer: B

Explanation:

A. **The de-duplication task module**

While de-duplication tasks can be accessed from task lists, this is not the preferred or recommended place to manage CMDB duplicates. It lacks the consolidated view, prioritization, and tooling needed for effective duplicate remediation.

B. ✓ The de-duplication dashboard on the CMDB Workspace

This is the preferred approach. The **CMDB Workspace de-duplication dashboard** provides a centralized, guided experience to identify, analyze, prioritize, and remediate duplicate CIs. It integrates with IRE-generated de-duplication tasks, supports bulk actions, templates, and best-practice workflows, making it the most efficient and governance-aligned option.

C. My Tasks in the Application Navigator

My Tasks shows assigned work across the platform but is not designed specifically for CMDB duplicate management. It does not provide the context, grouping, or remediation tools required for effective de-duplication.

Question: 89

A Configuration Management Governance team is transitioning from utilizing legacy CMDB status fields to CSDM life cycle status fields. Which table can be modified?

- A. Life Cycle Controls [life_cycle_control]
- B. Life Cycle Stage Status [life_cycle_stage_status]
- C. Life Cycle Stages [life_cycle_stage]
- D. Life Cycle Mapping [life_cycle_mapping]

Answer: D

Explanation:

A. Life Cycle Controls [life_cycle_control]

This table defines governance and enforcement logic for lifecycle behavior. It is system-controlled and not intended to be modified during a transition from legacy CMDB status fields.

B. Life Cycle Stage Status [life_cycle_stage_status]

This table contains the standard CSDM lifecycle stage statuses delivered out of the box. These values are part of the prescribed CSDM model and should not be customized.

C. Life Cycle Stages [life_cycle_stage]

This table defines the high-level lifecycle stages such as Plan, Build, Operate, and Retire. These stages are foundational to CSDM and should remain unchanged.

D. ✓ Life Cycle Mapping [life_cycle_mapping]

This is correct. The Life Cycle Mapping table is specifically designed to be configured during a transition from legacy CMDB status fields. It allows organizations to map existing legacy fields (such as install_status or operational_status) to the appropriate CSDM lifecycle stage and status without modifying the core CSDM tables.

Question: 90

What is the value of using the CMDB in security operations?

- A. Auto-resolves a vulnerability
- B. Allows security team to assess and remediate an incident
- C. Enables audits and attestations across CIs
- D. Identifies the IT infrastructure with a vulnerability

Answer: D

Explanation:

A. Auto-resolves a vulnerability

The CMDB does not automatically fix or resolve vulnerabilities. Remediation is handled by security tools, patching processes, or operational teams. The CMDB provides context, not automated resolution.

B. Allows security team to assess and remediate an incident

While the CMDB supports assessment by providing context, it does not by itself perform incident assessment or remediation. Those activities are carried out by Security Operations, Incident Response, and vulnerability management tools that use CMDB data.

C. Enables audits and attestations across CIs

Audits and attestations are governance and compliance activities. Although the CMDB supports them, this is not the primary value of the CMDB specifically within **security operations**.

D. ✓ Identifies the IT infrastructure with a vulnerability

This is correct. The key value of the CMDB in security operations is **providing visibility into which Configuration Items are affected by a vulnerability**. By linking vulnerabilities to servers, applications, services, and owners, the CMDB enables security teams to understand impact, prioritize risk, and route remediation to the correct teams.

Question: 91

A ServiceNow Administrator needs to create multiple new classes in the CMDB but wants to follow ServiceNow's best practices for naming CMDB tables to prevent technical debt. Which is the starting prefix for all custom CMDB tables?

- A. u_ci_cmdb
- B. ci_cmdb
- C. cmdb_ci
- D. u_cmdb_ci

Answer: D

Explanation:

D. u_cmdb_ci

Technical Breakdown

The u_ Logic: In the ServiceNow platform, the u_ prefix is the mandatory system-enforced identifier for all

custom tables created in the Global scope. This differentiates your organization's unique tables from those provided Out-of-the-Box (OOB) by ServiceNow.

The cmdb_ci Logic: This is the technical identifier for the **Configuration Item** base class. By including this in the table name, you ensure the table is properly categorized within the CMDB hierarchy.

The Full Sequence: When you create a custom class (e.g., for a "Security Appliance"), the resulting technical name is **u_cmdb_ci_security_appliance**.

Option-Wise Analysis

A. u_ci_cmdb (Incorrect):

This prefix order is non-standard. ServiceNow's internal logic always places the "cmdb" (the database identifier) before "ci" (the specific base class).

B. ci_cmdb (Incorrect):

This lacks the mandatory u_ for custom tables and follows the incorrect identifier order.

C. cmdb_ci (Incorrect):

This is the prefix for **standard Out-of-the-Box** tables. While it is the "exam-correct" prefix for OOB tables like cmdb_ci_server, it is not the prefix for custom tables. Using this would lead to potential naming conflicts during platform upgrades.

D. u_cmdb_ci (Correct):

This is the official and recommended prefix that satisfies both the "custom" requirement (u_) and the "CMDB hierarchy" requirement (cmdb_ci).

Question: 92

A CMDB Administrator has a report in ServiceNow that lists all CMDB Services that do not have an owner and wants to use a ServiceNow Playbook. What Governance process play can prevent this from recurring?

- A. Make the field Managed by mandatory on all CIs
- B. Set a default value on the Service Owner field so that it is never empty
- C. Make the field Owned by mandatory

Answer: C

Explanation:

A. Make the field Managed by mandatory on all CIs

The Managed by field represents the operational support group, not the business or service ownership. Making it mandatory does not ensure that services always have an owner defined.

B. Set a default value on the Service Owner field so that it is never empty

Setting a default value masks the governance issue. It can result in incorrect ownership data and does not enforce accountability or proper data stewardship.

C. ✓ Make the field Owned by mandatory

This is the correct governance control. Making the Owned by field mandatory ensures that every CMDB

Service must have a clearly defined owner at creation or update time, preventing ownerless services from being created in the future and aligning with CMDB governance best practices.

Question: 93

A CMDB Administrator has imported data into the ServiceNow CMDB from a third-party source using a Service Graph Connector. The Administrator wants to review specific field to field mappings for the import. Which feature will show that information?

- A. Integration Hub
- B. IntegrationHub ETL
- C. CMDB Integrations Dashboard

Answer: C

Explanation:

A. Integration Hub

Integration Hub is the platform used to build and run integrations and flows. It does not expose detailed, CMDB-specific field-to-field mappings created by Service Graph Connectors.

B. IntegrationHub ETL

IntegrationHub ETL focuses on transforming and loading data, but it is not the place to review Service Graph Connector mappings for CMDB ingestion.

C. ✓ CMDB Integrations Dashboard

This is the correct feature. The CMDB Integrations Dashboard provides visibility into Service Graph Connector activity, including **source-to-target (field-to-field) mappings**, ingestion status, and data quality details for CMDB imports.

Question: 94

A CMDB Administrator is comparing the Unified Map to the Service Mapping map. What are additional capabilities of the Unified Map?

- A. Map can be zoomed in and out
- B. Visibility to an application and the host it is installed on
- C. Number of levels displayed on a map can be modified
- D. Map nodes can be filtered based on user preferences

Answer: D

Explanation:

A. Map can be zoomed in and out

Zooming is a basic visualization capability and is available in both Unified Map and Service Mapping, so it is not an additional capability.

B. Visibility to an application and the host it is installed on

Service Mapping already provides this level of visibility, so this is not unique to the Unified Map.

C. Number of levels displayed on a map can be modified

Controlling the depth or levels of a map is also available in Service Mapping and therefore is not an additional capability.

D. ✓ Map nodes can be filtered based on user preferences

This is correct. The Unified Map adds more **user-driven interactivity**, allowing users to filter nodes based on preferences such as CI types or relationships, making it more flexible than traditional Service Mapping views.

Question: 95

A Configuration Manager is managing a CI class in the CMDB. The identification rule(s) needs an update. Where can the Configuration Manager view and configure the existing identification rule(s) for the class?

- A. IRE Application
- B. API Integrations
- C. CI Class Manager
- D. CI Identifiers module

Answer: C

Explanation:

C. CI Class Manager

Detailed Reasoning

The **CI Class Manager** is the centralized "Command Center" for everything related to a CI class. While the **CI Identifiers module** exists and shows a raw list of identifiers, ServiceNow officially recommends the CI Class Manager for administrative tasks for several reasons:

Unified View: It brings together **Identification Rules**, **Reconciliation Rules**, **Health Metrics**, and **Class Attributes** in one guided interface.

Hierarchy Awareness: It clearly shows whether a class is using its own unique rule or **inheriting** a rule from a parent class (e.g., how the Linux Server class might inherit the "Hardware" rule). This context is missing in the raw module view.

Inclusion Rules: Advanced logic like **Identification Inclusion Rules** (which you'll likely need for your Qualys project to filter out non-essential data) are also managed directly within the "Identification Rule" tab of the CI Class Manager.

Option-Wise Analysis

A. IRE Application (Incorrect):

The **Identification and Reconciliation Engine (IRE)** is a background process/framework, not an "application" with its own management UI.

B. API Integrations (Incorrect):

This refers to the method of sending data (REST, Import Sets), but not the configuration of the rules used once that data arrives.

C. CI Class Manager (Correct):

The official platform recommendation for viewing and configuring class-specific identification and reconciliation rules.

D. CI Identifiers module (Partially Correct/Distractor):

This module provides a simple list view of the cmdb_identifier table. While you can find the rules here, it is not the best-practice location for **managing a class** because it lacks hierarchy context.

Question: 96

The Configuration Management team finds value in the reports from CMDB 360/Multisource CMDB and wants to use it for all CI data. Which must be true in order for CMDB 360/Multisource CMDB to be able to report on and analyze that data?

- A. ServiceNow Discovery must be used to populate the CI data.
- B. The CI data must go through the IRE
- C. The CI data must be from an authorized Service Graph Connector.
- D. Reconciliation rules with priorities must be configured.

Answer: B

Explanation:

A. **ServiceNow Discovery must be used to populate the CI data**

CMDB 360 / Multisource CMDB does not require ServiceNow Discovery. CI data can come from many sources, including integrations, imports, and Service Graph Connectors.

B. ✓ **The CI data must go through the IRE**

This is required. CMDB 360 / Multisource CMDB relies on data being processed by the **Identification and Reconciliation Engine (IRE)** so it can track sources, reconcile attributes, analyze data quality, and report accurately across multiple data providers.

C. **The CI data must be from an authorized Service Graph Connector**

Service Graph Connectors are one way to ingest data, but they are not mandatory. Any data source is acceptable as long as it uses the IRE.

D. **Reconciliation rules with priorities must be configured**

Reconciliation rules improve data quality and conflict resolution, but they are not strictly required for CMDB 360 / Multisource CMDB reporting and analysis.

Question: 97

Where can a CMDB 360/Multisource CMDB Saved Query be viewed and created in the CMDB Workspace?

- A. Saved queries window on the CMDB 360 tab
- B. Saved queries window on the Insights tab
- C. Coverage window on the CMDB 360 tab
- D. CMDB Query Builder

Answer: A

Explanation:

A. ✓ **Saved queries window on the CMDB 360 tab**

This is the correct location. CMDB 360 / Multisource CMDB saved queries are **created and managed directly within the CMDB 360 tab** in the CMDB Workspace. These saved queries are used to analyze CI data across multiple sources and are specific to CMDB 360 reporting.

B. **Saved queries window on the Insights tab**

The Insights tab focuses on health metrics, dashboards, and analytics. It does not provide the interface for creating or managing CMDB 360 saved queries.

C. **Coverage window on the CMDB 360 tab**

The Coverage window is used to understand data source coverage across CIs and attributes. It is not where saved queries are created or viewed.

D. **CMDB Query Builder**

CMDB Query Builder is a general-purpose query tool and is not specific to CMDB 360 / Multisource CMDB saved queries.

Question: 98

A CMDB Administrator is beginning the journey of populating the CMDB and needs to verify that any data which is no longer useful/applicable is removed. Which governance management tool will accomplish this?

A. CMDB and CSDM Data Foundations Dashboard

B. CMDB Data Manager

C. CMDB Health Dashboard

D. CI Class Manager

E. De-duplication Templates

Answer: B

Explanation:

A. **CMDB and CSDM Data Foundations Dashboard**

This dashboard provides visibility, metrics, and playbooks for CMDB health and governance, but it does not directly remove or retire obsolete data.

B. ✓ **CMDB Data Manager**

This is the correct tool. CMDB Data Manager is used for **governance and lifecycle management of CI data**, including verification, attestation, certification, and **retirement or removal of data that is no longer useful or applicable**.

C. **CMDB Health Dashboard**

The Health Dashboard measures completeness, correctness, and compliance, but it does not govern lifecycle removal of CIs.

D. **CI Class Manager**

CI Class Manager manages class structure, attributes, and rules, not data lifecycle cleanup.

E. **De-duplication Templates**

De-duplication templates address duplicate CIs, not obsolete or no-longer-applicable data.

Question: 99

User endpoint devices are imported into the CMDB and populate the 'Assigned to' [assigned_to] field on the Computer [cmdb_ci_computer] CI. The Asset team puts in a request for the Configuration Analysts to populate the 'Assigned to' field on the related Asset. What action does a Configuration Analyst take to achieve this in an automated way?

- A. Use the Asset-CI Field Mapping module to create a new rule to replicate the 'Assigned to' value between the asset and associated CI
- B. Hide the 'Assigned to' field on the asset record and create a new field that dot walks to the related CI to get the 'Assigned to' value
- C. Configure a business rule on the computer table to use a script to populate the 'Assigned to' field on the asset based on insert or update in the computer class 'Assigned to' field

Answer: A

Explanation:

A. ✓ **Use the Asset-CI Field Mapping module**

This is the recommended and fully supported approach. The Asset-CI Field Mapping module is designed specifically to keep asset and CI fields in sync. Creating a mapping rule allows the Assigned to value on the Computer CI to be automatically populated on the related Asset record without custom code, aligning with ServiceNow best practices.

B. **Hide the field and dot-walk to the CI**

This does not actually populate the asset record. It only displays the value and can cause reporting, integrations, and downstream processes to fail because the asset field remains empty.

C. **Create a custom business rule**

While technically possible, this introduces custom scripting, increases technical debt, and bypasses the standard, supported synchronization mechanism provided by ServiceNow.

Question: 100

Choose 2 option. What are the characteristics or functions of ServiceNow IntegrationHub ETL?

- A. Performs discovery data collection and updates the CMDB
- B. Integrates third-party data into the CMDB or into non-CMDB tables
- C. Uses the IRE to process and integrate data
- D. Imports Microsoft SCCM/Intune data into the CMDB

Answer: B,C

Explanation:

A. Performs discovery data collection and updates the CMDB

This describes **ServiceNow Discovery**, not IntegrationHub ETL. ETL does not actively discover infrastructure.

B. ✓ Integrates third-party data into the CMDB or into non-CMDB tables

This is a core function of IntegrationHub ETL. It extracts data from external sources, transforms it, and loads it into either CMDB tables or non-CMDB tables.

C. ✓ Uses the IRE to process and integrate data

When IntegrationHub ETL loads data into the CMDB, it leverages the **Identification and Reconciliation Engine (IRE)** to correctly identify CIs, prevent duplicates, and reconcile attributes.

D. Imports Microsoft SCCM/Intune data into the CMDB

While technically possible, this is **not the intended or fastest approach**. ServiceNow best practice for SCCM or Intune ingestion is **Service Graph Connectors**, not IntegrationHub ETL.

Question: 101

A CMDB Administrator wants to educate the team on the various actions that can be performed within the CMDB Workspace. What actions can be initiated from the CMDB Workspace?

- A. Create a new CMDB class
- B. Execute ServiceNow Discovery
- C. Remediate duplicate CI records
- D. Create a CMDB Data Manager certification policy

Answer: C

Explanation:

A. Create a new CMDB class

CMDB classes are created and managed in **CI Class Manager**, not from the CMDB Workspace.

B. Execute ServiceNow Discovery

Discovery is configured and executed from the **Discovery application**, not initiated from the CMDB Workspace.

C. ✓ Remediate duplicate CI records

This is correct. The CMDB Workspace provides **de-duplication dashboards and tools** that allow administrators to review and remediate duplicate CI records generated by the IRE.

D. Create a CMDB Data Manager certification policy

Certification policies are created and managed in **CMDB Data Manager**, not directly from the CMDB Workspace.

Question: 102

A CMDB Administrator wants to ensure that only relevant CIs from managed classes will be shown on Incident, Problem, and Change records. Which checkbox needs to be checked in the CI Class Manager for the CMDB Administrator to achieve the requested result?

- A. Principal Class
- B. Main Record
- C. Independent
- D. Extensible

Answer: A

Explanation:

A. ✓ Principal Class

Checking the Principal Class checkbox ensures that only CIs from relevant, managed classes are available for selection on **Incident, Problem, and Change** records. This helps limit CI selection to meaningful classes and improves data quality and usability for ITSM processes.

B. Main Record

This controls how records are displayed or treated internally, not whether CIs appear in task records like Incident or Change.

C. Independent

Independent classes do not inherit from other classes, but this setting does not control CI visibility on task forms.

D. Extensible

Extensible determines whether a class can be extended by child classes. It has no impact on CI selection in Incident, Problem, or Change forms.

Question: 103

A CMDB Administrator wants to ensure all short-lived CIs that have not been discovered in the past week are removed. After retiring the CI records, which recommended action does the CMDB Administrator take?

- A. Create a delete policy
- B. Create a scheduled job
- C. Create a business rule

Answer: A

Explanation:

A. ✓ Create a delete policy

This is the recommended action. After CIs are **retired**, a **CMDB Data Manager delete policy** is used to automatically and safely remove those retired records after defined conditions are met. This aligns with CMDB lifecycle governance and avoids manual cleanup.

B. Create a scheduled job

Scheduled jobs are generic automation tools and are not the recommended governance mechanism for controlled CI deletion.

C. **Create a business rule**

Business rules introduce custom logic and technical debt. ServiceNow recommends using **Data Manager policies** instead of scripting for lifecycle cleanup.

Question: 104

A CMDB Administrator needs to set a CI Class as a Principal Class. Which CI Class Manager tab would need to be accessed?

- A. Class Info > Basic Info
- B. Health > Attributes
- C. Class Info > Attributes

Answer: A

Explanation:

A. ✓ **Class Info > Basic Info**

The **Principal Class** checkbox is configured on the **Basic Info** sub-tab under **Class Info** in the CI Class Manager. This is where high-level class behaviors and flags, such as Principal Class, are defined.

B. **Health > Attributes**

This area is used for configuring health-related settings and attributes, not class visibility or principal status.

C. **Class Info > Attributes**

This tab manages class attributes and fields, not class-level flags like Principal Class.

Question: 105

A CMDB Administrator changes the query for the SCCM Service Graph Connector. What is the impact of this change?

- A. Any updates for the SCCM Service Graph Connector will be skipped during the upgrade.
- B. Any Scheduled Jobs for the SCCM Service Graph Connector will need to be configured.
- C. The Data Source for the SCCM Service Graph Connector will be marked as Inactive.

Answer: C

Explanation:

When a CMDB Administrator **changes the query** used by a **Service Graph Connector** (such as SCCM), ServiceNow treats this as a deviation from the certified, supported configuration. To protect CMDB data quality and connector integrity, **ServiceNow automatically marks the connector's Data Source as Inactive**. This prevents the connector from continuing to ingest data with a modified query that may no longer align with the supported data model.

Why the other options are incorrect:

Option A is incorrect because modifying the query does not merely skip future updates. Service Graph Connectors are not upgrade-skipped in the same way as traditional customizations; instead, the connector is disabled by inactivating the data source.

Option B is incorrect because scheduled jobs are already part of the Service Graph Connector setup. Changing the query does not require reconfiguring schedules.