ServiceNow CIS
Data Foundations Guide

CMDB & CSDM

A practical quick-revision guide for CIS Data Foundations: focused concepts, exam cues, examples, and a consolidated practice bank.

Attia Virk | ServiceNow Developer & Consultant

What this guide is A quick-revision companion built around the current CIS-DF blueprint and official ServiceNow documentation. It does not replace the required training; use it after the learning path to consolidate the concepts most likely to be tested.

Publishing note The practice questions are original/paraphrased study questions inspired by common learning objectives and practice scenarios. They are not presented as verbatim live-exam content.

1. Start Here: Your Preparation Roadmap

Use the official training as your foundation, then use focused revision and practice questions to test whether you can apply the concepts.

Step	What to do	Mentor guidance
1	Confirm eligibility	Open the CIS - Data Foundations credential in ServiceNow University and confirm your current eligibility, attempt policy, and scheduling options.
2	Complete the prerequisite training	CMDB Fundamentals is the key eligibility course. Complete the labs, not just the videos.
3	Finish recommended learning	Work through the Data Foundations Implementer path and current CSDM/CMDB learning content shown in your learning plan.
4	Read the blueprint before revising	Spend study time according to the weights. Govern is the largest domain.
5	Use a lab/PDI where available	Open CI Class Manager, CMDB Workspace, Data Manager, CMDB Health, Query Builder, and Unified Map.
6	Do the quick revision	Focus on distinctions rather than memorizing isolated facts.
7	Attempt the practice bank	Answer without looking at green options; then read the explanation.
8	Final review	Revisit weak blueprint domains and the conflict-traps page rather than random answer keys.

Mentor note Practice questions should not be your first source of learning. The exam rewards understanding why ServiceNow prefers one native, governed capability over another.

Recommended official starting points

- CIS - Data Foundations credential / blueprint
- CIS-DF eligibility badge / CMDB Fundamentals
- Data Foundations Implementer learning path

2. Exam Blueprint: Study in the Same Shape as the Exam

The current blueprint is organized into five learning domains. Use the percentages as a study-time guide rather than treating every topic equally.

Blueprint domain	Weight	What to revise
Configuration	15%	CI Class Manager, hierarchy/attributes, IRE, identification/reconciliation, CMDB 360
Ingest	19%	Discovery, ACC, Service Graph Connectors, IntegrationHub ETL, Import Sets, relationships, Asset-CI alignment
Govern	35%	CMDB Health, Data Manager, duplicates, roles, Principal Classes, lifecycle, operationalization
Insight	20%	Query Builder, Unified Map, CMDB 360 saved queries, Data Foundations dashboards/playbooks, business value
CSDM Fundamentals	11%	CSDM domains, stages, stakeholders, key objects/relationships and outcomes

Study-time shortcut If you have 10 hours left: roughly 3.5 hours Govern, 2 Insight, 2 Ingest, 1.5 Configuration, and 1 CSDM - then rebalance for your weak areas.

1. Configuration | 15%

Configuration is about the structure and rules that make CMDB data trustworthy before you start loading large volumes of CIs. For the exam, focus on where classes are defined, how IRE decides whether a CI already exists, and which source is allowed to update which attributes.

CI Class Manager - your control room for the class model

Concept	What it means	Simple example	What to remember for MCQs
Class hierarchy	CMDB classes inherit fields and behavior from parent classes. A more specific class should sit under the correct parent.	Windows Server is a child of Server.	Parent -> child means more specific. Reuse an existing class before creating a custom one.
Class attributes	Fields can be inherited from a parent or introduced on the selected class.	Serial number may exist higher in the hardware hierarchy; a Windows-specific field may be added lower.	Know the difference between inherited attributes and attributes added on the selected class.
Principal Class	Marks managed classes that should be favored in CI reference experiences so users do not browse the entire CMDB hierarchy.	Incident users should select meaningful managed CI classes rather than every technical sub-class.	Principal Class is class configuration, not a CMDB Health metric.
Identification rules	Define the attributes IRE uses to decide whether an incoming CI matches an existing CI.	A server might be identified by serial number, then by other identifier entries in priority order.	Identification answers: 'Is this the same CI?'
Reconciliation rules	Control which authorized source may update a class or specific attributes when multiple sources report the same CI.	Discovery may own OS version while an asset source owns warranty data.	Reconciliation answers: 'Who is allowed to update it?'

Exam shortcut: If the question asks where to inspect class hierarchy, attributes, identifiers, or Principal Class settings, think CI Class Manager.

IRE - Identification and Reconciliation Engine

IRE is the gatekeeper for CI identity and source authority. A useful mental model is: first identify the record, then reconcile the incoming values. This is why direct writes to CMDB tables are risky when the integration should be governed by IRE.

Step	Question IRE answers	Example
1. Identify	Does this incoming payload represent an existing CI or a new CI?	Incoming Windows server matches an existing CI by an identification rule.
2. Reconcile	If it is the same CI, may this source update these attributes?	SCCM can update software details, but a higher-priority source may own another field.
3. Update / Insert	Update the authorized existing CI, or insert a new CI when no match exists.	The matched server is updated instead of creating a duplicate.

Common trap: A lower-priority source can still be useful. Reconciliation limits updates to existing records; it does not automatically mean that source can never insert a new CI.

Source precedence, Data Refresh Rules, and CMDB 360

Feature	Purpose	Example / exam cue
Static reconciliation priority	Ranks sources for a class or attribute.	A source with stronger authority wins updates while it remains current.
Data Refresh Rule	Lets a lower-priority source update after a higher-priority source has been inactive for a defined time.	Primary connector stops reporting for several days; secondary source is allowed to refresh the CI.
CMDB 360 / Multisource CMDB	Retains and exposes source-specific values for the same CI so you can analyze how sources disagree.	Compare an attribute reported by Discovery and SCCM.
Dynamic reconciliation	Uses retained multisource values to choose a value dynamically instead of only using a fixed source priority.	Choose the most recently reported or largest value when that rule is appropriate.

Reclassification - Upgrade, Downgrade, Switch

Operation	Direction	Example	Risk
Upgrade	Parent -> child (more specific)	Server -> Windows Server	Normally no class-specific data is lost because the CI is moving to a more specific child.
Downgrade	Child -> parent (less specific)	Windows Server -> Server	Child-only attributes may no longer fit; data loss is possible.
Switch	Sibling -> sibling	Linux Server -> Windows Server	Class-specific attributes can be excluded; data loss is possible.

Memory anchor: UP = more specific. DOWN = less specific. SWITCH = same level, different sibling.

2. Ingest | 19%

Ingest questions test whether you choose the cleanest supported way to populate CMDB and preserve IRE behavior. The exam often gives several technically possible options; prefer the native, governed integration path with the least custom debt.

Choose the right ingestion method

Method	What it is	Best-fit example	Exam clue
ServiceNow Discovery	Agentless infrastructure discovery using MID Server and patterns.	Discover servers and network devices from IP ranges.	Think infrastructure, patterns, automated technical relationships.
Agent Client Collector (ACC)	Agent-based collection from managed endpoints.	Collect endpoint telemetry when agents are preferred or network scanning is limited.	Agent-based = ACC.
Service Graph Connector	Packaged connector for supported third-party sources.	Bring SCCM or another supported vendor dataset into CMDB.	Fastest time to value when a supported connector exists.
IntegrationHub ETL	Low-code ETL framework for transforming external data into ServiceNow data models, including CMDB with IRE-aware patterns.	Map third-party asset or infrastructure data where a packaged connector is not the right fit.	Prefer supported ETL mapping over direct table writes.
Import Set + Transform Map	Staging-table import approach; for CMDB, route records through IRE instead of ordinary coalesce-only logic.	Custom CSV/API load into a staging table.	CMDBTransformUtil is used in the transform approach to invoke IRE.

Decision rule: Supported connector available? Prefer it. No connector? Use a supported IRE-aware integration pattern. Avoid a custom Business Rule whose only purpose is to force CMDB synchronization.

Service Graph Connector - what to know

Topic	One-line explanation	Example
Field mappings	Translate vendor/source fields into the ServiceNow target data model.	SCCM device fields map to Windows Server attributes.
IRE processing	The connector should create/update CIs through the platform's identification and reconciliation framework.	Existing CI is updated instead of blindly inserting another row.
Customization risk	Changing baseline connector mappings increases upgrade and support responsibility.	A customized transform may be skipped during an upgrade and require review.
Service Graph Connector Central	Central experience for administering supported Service Graph Connectors.	Check connector configuration, runs, and mapping-related setup.

Asset-CI synchronization and non-discoverable data

Scenario	Preferred thinking	Example
Asset and CI share fields	Use the platform's supported Asset-CI field synchronization/mapping before custom scripting.	Assigned to is synchronized between paired Asset and CI records.
Discovery cannot collect a business attribute	Use the authoritative external source and let integration/reconciliation govern updates.	Warranty expiration comes from the asset/ERP source, not from Discovery.
Relationship creation	Use Discovery/Service Mapping where the topology can be discovered; do not manually maintain thousands of technical relationships.	Application Runs on Server relationships are created from discovered topology.

Service Mapping / application-service creation patterns

Approach	What it does	Easy example
Top-down Service Mapping	Starts from an application/service entry point and discovers dependencies using patterns.	Map a mission-critical customer portal from its entry point through web/app/database tiers.
Connection Suggestions	Uses observed connections to suggest service membership; useful in dynamic cloud/container environments.	Suggest resources that communicate with known application components.
Tag-based	Uses cloud/resource tags to define application-service membership.	All resources tagged App=Billing become members of the Billing service.
Dynamic CI Group	Groups CIs by saved query or criteria rather than discovering an end-to-end dependency path.	All production Windows servers supported by one team.

3. Govern | 35%

Govern is the largest exam domain. The key theme is not 'clean the CMDB once'; it is continuous control: measure data quality, assign ownership, create remediation work, and manage CI lifecycle through policies.

CMDB Health - understand the scorecards, not just their names

Scorecard	What it asks	Typical metrics	Example
Completeness	Do the CIs contain the data we expect?	Required fields, Recommended fields	A Server is missing Support Group.
Correctness	Is the CI data trustworthy and current?	Duplicates, Orphans, Staleness	Two records represent the same server, or a CI has not been updated recently.
Compliance	Does the CI match a defined expected configuration?	Desired State / scripted audits	A server should have an approved configuration value but fails the audit.

Memory anchor: Completeness = missing data. Correctness = bad/duplicate/stale data. Compliance = expected state.

Required vs Recommended fields

Required	Recommended
Globally enforced / mandatory where configured. Missing values can block valid record handling or cause integrations to fail.	Not globally mandatory. Used as a quality target and a safer proving ground before making a field required.
Use sparingly and only for data that truly must always exist.	Good for fields such as ownership/support information where you want visibility and remediation tasks.
Exam idea: 'must be populated' points to Required.	Exam idea: 'best-practice field missing' points to Recommended.

Staleness, inclusion rules, and task creation

Concept	What it means	Example / exam clue
Staleness	Flags CIs that have not been updated within the configured effective duration.	Default study value: 60 days; evaluation is based on Updated (sys_updated_on).
Health Inclusion Rule	Narrows which CIs/classes are evaluated by a health metric.	Orphan evaluates 4,200 CIs while Duplicate evaluates 5,000 because Orphan has a narrower inclusion rule.
Health task creation	A health metric can be configured to create tasks for failed CIs and route them to an assignee group.	Recommended-field failures create work for the responsible data owner.

Duplicate management - detect, choose the main CI, preserve good data

Duplicate remediation is more than deleting the newer or older record. The goal is to converge on one trusted CI while preserving authoritative attributes and valid relationships.

Capability	Role in duplicate management
CMDB Health - Duplicate metric	Highlights duplicate problems as part of Correctness.
De-duplication tasks / Workspace	Provides the operational queue and wizard-driven remediation experience.
Duplicate CI Remediator	Helps select the main CI and merge useful attributes/relationships from duplicate records.
De-duplication templates	Support repeatable/bulk duplicate remediation scenarios.

Exam scenario: If the discovered CI has the freshest IP but the manually loaded duplicate has a critical Business Application relationship, keep the best main CI and preserve/merge the valuable relationship - do not blindly delete it.

CMDB Data Manager - lifecycle governance

Policy type	Purpose	Example
Retire	Moves CIs that are no longer active toward a retired lifecycle state.	A server has been decommissioned and should no longer be treated as operational.
Archive	Moves inactive/retired records out of the active working set while retaining them for historical needs.	Old retired CIs are retained but removed from day-to-day CMDB volume.
Delete	Permanently removes records when governance criteria are met.	Short-lived retired CIs are deleted after the defined retention logic.
Attestation	Asks responsible users to confirm that the CI still exists / is still valid.	A data-center owner confirms each server physically or logically still exists.
Certification	Asks responsible users to validate specified CI attributes and, when allowed, correct them.	A service owner verifies ownership and support fields.

Do not mix these: De-duplication and reclassification are important CMDB capabilities, but they are not Data Manager policy types.

Remediation and Data Foundations playbooks

Item	What to remember
CMDB Remediation Rule	Connects a failed health condition to remediation logic/workflow and defines how remediation work is triggered.
Data Foundations dashboard	Surfaces indicators of data quality and foundational gaps so teams can prioritize remediation.
Get Well / remediation playbooks	Explain the problem, likely impact, and recommended remediation steps. They guide work; they are not magic one-click fixes.

4. Insight | 20%

Insight is about using the CMDB, not just maintaining it. Questions here usually ask where you analyze topology, build relationship queries, compare multiple sources, or use dashboard metrics to understand operational impact.

CMDB Workspace and Unified Map

Feature	What it gives you	Example
Unified Map	Visual topology around a selected CI/service, with controls to filter what is shown.	View a service, its dependent servers, and related operational context.
Node filters	Reduce visual noise using filters such as CI type/class or discovery source.	Hide non-server nodes or show only CIs from a particular discovery source.
Related items/context	Shows operational records related to the selected CI.	Check active incidents or problems for a CI without leaving the topology view.
Natural Language Query	Lets an administrator ask CMDB questions in natural language where the feature is available in CMDB Workspace.	Ask for operational Windows servers in a location instead of manually building every condition.

CMDB Query Builder

Query Builder is about graph thinking: add the CI classes you need, connect them with relationships, filter the node that owns the attribute, then optionally bring in non-CMDB tables for operational context.

Need	How to think about it
Only Seattle databases	Put Location = Seattle on the Database node, because that attribute belongs to the database CI.
Databases connected to Application Services	Add both classes and define the relationship path between them.
Also include Incidents	Add Incident from the non-CMDB tables list and relate it to the CI context.
Show fields as columns	Columns control the output display; filters control which records qualify. Do not confuse the two.

Exam trap: A column is not a filter. If the question says 'only CIs where X = Y', configure a filter on the node that owns X.

CMDB 360 / Multisource analysis

Concept	Meaning
Saved queries	Can be created/viewed from the CMDB 360 experience in CMDB Workspace.
Coverage / source comparison	Shows how multiple data sources contribute values for the same CI.
Why it matters	Helps diagnose conflicting data, understand source coverage, and support dynamic reconciliation decisions.

Business value of trusted CMDB data

Process	Value from CMDB
Incident / Problem	Faster context, affected-CI visibility, ownership routing, and dependency understanding.
Change	Impact analysis and better routing/assignment using CI and service ownership context.
Event Management	Bind alerts to specific CIs and understand the business/service impact of technical events.
Security Operations	Identify vulnerable infrastructure and add CI/business context for risk prioritization and remediation.
Operational resilience	Understand dependencies and critical services so teams can assess outage impact more quickly.

5. CSDM Fundamentals | 11%

CSDM is the prescriptive data model for placing service-related records in the right tables and connecting them with the right relationships. The exam tests both where an object belongs and how an organization matures through Foundation -> Crawl -> Walk -> Run -> Fly.

First distinction: CSDM domains vs implementation stages

Domains	Implementation stages
Describe where data belongs in the conceptual model - for example Design & Planning, Service Delivery, Service Consumption.	Describe how far your organization has progressed in adopting CSDM - Foundation, Crawl, Walk, Run, Fly.
Question sounds like: 'Which domain does this object belong to?'	Question sounds like: 'At which stage do we start doing this?'

Foundation -> Crawl -> Walk -> Run -> Fly

Stage	What it means in plain English	Quick example	MCQ memory
Foundation	Prepare clean referential data and base information that other CSDM records depend on.	Standardize foundational records such as locations, companies, users/groups and other reference data needed for reliable reporting.	Get the foundation right before modeling services.
Crawl	Work with core CMDB data that supports ITSM and establish a reliable operational CI base.	Incidents and changes can point to meaningful CIs with proper class and ownership context.	Crawl = core ITSM/CMDB usage.
Walk	Identify the infrastructure and applications technical teams support and organize them with technology management services/offerings and group context.	A Windows team owns a Technology Management Offering that manages a group of production Windows servers.	Walk = technical support model + managed infrastructure.
Run	Connect technology delivery to the business that consumes and sells/uses the service.	Application services and technical offerings are related to business-facing services/offerings.	Run = technology meets business.

Stage	What it means in plain English	Quick example	MCQ memory
Fly	Complete or nearly complete the CSDM model and use it for broader portfolio, service, data, and optimization use cases.	Use Business Capabilities, Information Objects, Business Services and portfolio analysis for investment/risk decisions.	Fly = mature, end-to-end model.

Important nuance: Information Objects are often associated with the mature/Fly model, but ServiceNow explicitly notes they can be implemented earlier when business requirements need them.

High-yield CSDM objects

Object	What it represents	Easy example	Where to think of it
Business Application	A design/portfolio view of an application - what the organization owns, invests in, rationalizes, and governs.	Payroll Application as an application portfolio record.	Design & Planning.
Application Service / Service Instance	The operational realization/deployment of an application or service that can be mapped to supporting CIs.	Payroll-PROD running on web/app/database infrastructure.	Service Delivery.
Technology Management Service	A technical service that represents technology delivered and managed by technical teams.	Windows Hosting service.	Service Delivery / technical management.
Technology Management Offering	A specific offering of a technical service, often carrying environment/location/support group context.	Windows Hosting - Production.	Walk-stage technical support context.
Business Service	A business-facing service that consumers recognize.	Employee Payroll Service.	Service Consumption.
Business Service Offering	A consumable variation/package of a Business Service.	Payroll Service - North America.	Service Consumption.
Business Capability	A high-level business ability, independent of a specific application.	Payroll Management capability.	Design / portfolio planning context.

Information Objects - what they actually are

An Information Object represents the logical type of data used by a Business Application. It is not the database server and not a random text label. It helps you describe what information an application creates, reads, updates, or deletes and is especially useful for risk, privacy, and architecture questions.

Question	Answer
What does it represent?	Logical information/data, such as Customer PII, Payment Card Data, Patient Health Information, or Employee Records.
What is it related to?	Business Applications. The relationship lets you understand which applications use which types of information.
Why does the exam care?	Because it connects application architecture with data/risk context - especially PII, PCI, HIPAA, and other regulated information.
Simple example	A Hospital Scheduling Business Application uses an Information Object called Patient Health Information. Security/Risk teams can then identify which applications handle that sensitive information.
What should I not confuse it with?	A Database Instance or Database Catalog is the physical/technical data-store side; an Information Object is the logical information being handled.

MCQ clue: If the scenario says 'what type of sensitive data does this Business Application use?' - Information Object should be in your head immediately.

CSDM exam mental model

If the question says...	Think...
Portfolio, investment, rationalization of an application	Business Application
Operational deployment / topology / service map	Application Service / Service Instance
Technical team supports infrastructure as an offering	Technology Management Service / Offering
Business users consume a service	Business Service / Business Service Offering
Sensitive or regulated data used by applications	Information Object
What the business is capable of doing	Business Capability

PRACTICE QUESTION BANK

150+ scenario, MCQ, and matching questions

3. Practice Question Bank

Use this as one continuous practice bank. Exact duplicates have been removed, but wording and scenario variations are intentionally retained so you can recognize the same concept when it is tested differently. Correct options are highlighted in green; each question includes a short explanation. Drag-and-drop/matching questions are included in the same continuous numbering.

Q1. Where are CI identification rules configured?

A. CMDB Health Dashboard

B. CMDB Workspace home only

C. CMDB Data Manager

D. CI Class Manager

Why: Identification rules are managed from CI Class Manager for the relevant class.

Q2. A child CI class should fall back to its parent class identification rules when no child rule matches. Which setting enables this behavior?

A. Allow fallback to parent's rules

B. Independent

C. Applies to

D. Criterion attributes

Why: The fallback option allows parent identification behavior to be used.

Q3. An Apache Web Server CI is rediscovered after its version changes. What should happen if identification matches the existing CI and the source is authorized?

A. The CI is duplicated

B. The existing CI is reconciled and the version is updated

C. The CI is always reclassified

D. A new CI is always inserted

Why: IRE identifies the existing CI and reconciliation determines whether the incoming source can update the attribute.

Q4. A server record is reclassified from Server to Windows Server or Linux Server during Discovery. What kind of operation is this?

A. Class Upgrade

B. Class Switch

C. Class Downgrade

Why: Moving from a more generic parent class to a more specific child class is a class upgrade.

Q5. Several authorized data sources can update the same class and attributes. How should the authoritative source be controlled?

A. Manually run sources in order

B. Use identification rule run order

C. Use reconciliation priorities

D. Use a UI policy

Why: Reconciliation rules establish source precedence for attribute updates.

Q6. With reconciliation priorities ServiceNow=100, Altiris=200, SCCM=300, which general principle is correct?

A. A lower numeric priority represents higher precedence

B. A higher number always overrides a lower number

C. Priority only affects inserts

D. Reconciliation does not consider discovery source

Why: Reconciliation priority is used to determine which source is more authoritative; lower priority numbers take precedence.

Q7. Can a lower-authority source still insert a brand-new CI when no matching CI exists?

A. No

B. Yes, if identification finds no existing match and the source is allowed to create the CI

Why: Reconciliation primarily governs attribute update authority after identification; an unmatched payload can still create a CI when permitted.

Q8. A Platform Data Owner has several discovery sources but does not use CMDB 360. Can standard IRE reconciliation rules still govern source precedence?

A. No, CMDB 360 is mandatory for reconciliation

B. Yes, IRE reconciliation is a core CMDB capability independent of the CMDB 360 visualization feature

Why: IRE identification/reconciliation logic is not dependent on using CMDB 360 as the interface for source history.

Q9. What is the purpose of dynamic reconciliation rules such as Most Recently Reported or Largest Value?

A. Choose attribute values dynamically according to rule logic instead of fixed source priority alone

B. Delete duplicate CIs

C. Create CI classes

D. Replace identification rules

Why: Dynamic reconciliation selects an attribute value using value/time-based logic across reported sources.

Q10. A new IoT Sensor class is needed. Which two practices are recommended?

A. Delete an unused class and reuse it

B. Check/update the CMDB CI Class Models Store app to see whether the class already exists

C. Add the new class under the most appropriate parent

D. Modify an unrelated existing class

Why: Reuse the supported class model where possible; if a new class is truly required, extend the correct parent rather than repurposing unrelated classes.

Q11. A new custom CI class is needed. Which minimum roles were identified in the study material for creating the class?

A. sn_cmdb_admin and personalize_dictionary

B. cmdb_inst_admin and personalize_form

C. data_classification_admin and personalize_dictionary

D. itil_admin and personalize_form

Why: Class creation requires CMDB administration plus dictionary-level customization capability.

Q12. Which tab in CI Class Manager shows attributes introduced specifically on the selected class rather than inherited from a parent?

A. Added

B. Child

C. All

D. Derived

Why: The Added view isolates attributes defined on that class.

Q13. Where are Principal Classes designated?

A. CI Class Manager

B. System Properties only

C. CMDB Data Manager

D. CMDB Workspace home page only

Why: Principal class configuration is managed in the class-management experience.

Q14. A CI Class Owner wants quick access to a frequently managed class such as Windows Server. Which feature helps?

A. Pinned Classes

B. Search CI Classes only

C. CI Favorites in a separate app

Why: Pinned Classes provide quick access to commonly managed classes.

Q15. A Configuration Manager wants to inspect how CIs are connected in CMDB Workspace. Which feature should be used?

A. Unified Map

B. Business Service Map only

C. Relationship Map outside Workspace only

Why: Unified Map visualizes relationships and contextual information around CIs/services in CMDB Workspace.

Q16. Which system property enables CMDB 360 / multisource functionality?

A. glide.identification_engine.multisource_enabled

B. glide.ui.multisource_enabled

C. glide.cmdb.legacy_multisource

Why: This property enables multisource functionality for IRE/CMDB 360.

Q17. Which property enables capturing multisource data for CIs in CMDB classes?

A. glide.identification_engine.multisource_cmdb_ci_enabled

B. glide.cmdb.capture_all

C. glide.discovery.multisource_ci

Why: This property controls CI-level multisource data capture.

Q18. What are three common ways total cost of ownership information for Business Applications can be brought into ServiceNow?

A. API/integration ingestion

B. Manual entry or data load

C. Scheduled integration/job

D. Treating an ad hoc spreadsheet as automatic governance by itself

Why: TCO data can be integrated, loaded/entered, or refreshed on a schedule; the key is having an owned and governed source.

Q19. Which plugin/license component must be active to enable CMDB 360?

A. ITOM Discovery License plugin

B. HR Service Delivery

C. Software Asset Management Professional

D. Virtual Agent

Why: Current CMDB 360 setup instructions require the ITOM Discovery License plugin.

Q20. A team uses ordinary IRE reconciliation rules but has not enabled CMDB 360. Which statement is most accurate?

A. IRE cannot reconcile anything without CMDB 360

B. Standard IRE reconciliation still controls authorized source updates, but CMDB 360-specific source history/dynamic reconciliation is unavailable

C. All sources become equal priority

D. Identification rules stop working

Why: IRE is the core identification/reconciliation framework. CMDB 360 adds multisource retention/visualization and enables dynamic reconciliation.

Q21. Which feature must be enabled before using Dynamic Reconciliation rules such as Largest Value or Most Reported?

A. CMDB 360

B. Data Certification

C. Agent Client Collector

D. Service Portal

Why: Dynamic reconciliation evaluates values retained by CMDB 360 across discovery sources.

Q22. A higher-priority discovery source has stopped reporting. Which feature can allow a lower-priority source to update the CI after a defined period?

A. Health Inclusion Rule

B. Data Refresh Rule

C. Identification Rule

D. Principal Class Filter

Why: Data Refresh Rules work with static reconciliation rules to determine when a lower-priority source may update after the authoritative source has not refreshed the data.

Q23. Which statement correctly distinguishes static and dynamic reconciliation?

A. Static rules use source priority; dynamic rules can choose values such as largest or most reported across CMDB 360 data.

B. Dynamic rules identify new CIs while static rules delete duplicates.

C. Static rules require CMDB 360 but dynamic rules do not.

D. They are identical.

Why: Static reconciliation is source-priority based. Dynamic reconciliation evaluates multisource values and therefore requires CMDB 360.

Q24. A Windows Server CI is reclassified to its parent Server class. What reclassification type is this?

A. Upgrade

B. Downgrade

C. Switch

D. Promotion

Why: Moving from a more specific child class to a less specific parent class is a class downgrade.

Q25. A Linux Server CI is reclassified to Windows Server, where both are siblings under Server. What type of reclassification is this?

A. Upgrade

B. Downgrade

C. Switch

D. Merge

Why: Moving between sibling classes is a class switch.

Q26. What is the main purpose of Principal Classes?

A. Focus operational CI reference choices on classes the organization manages

B. Make those classes bypass IRE

C. Disable Discovery for other classes

D. Turn classes into Asset tables

Why: Principal Classes help focus operational use and reference fields on the CI classes that matter to the organization.

Q27. When Import Sets and Transform Maps are used to load CI data, what Transform Map script timing is used to call IRE via CMDBTransformUtil?

A. onStart

B. onBefore

C. onComplete

D. onAfter

Why: ServiceNow documents use of CMDBTransformUtil in an onBefore Transform Map script to route rows through IRE.

Q28. Which API is commonly used in a Transform Map to process CMDB rows through IRE and minimize duplicate CIs?

A. GlideAggregate

B. CMDBTransformUtil

C. GlideSchedule

D. SNC.AssetManagement

Why: CMDBTransformUtil invokes Identification and Reconciliation for Import Set rows.

Q29. What is the recommended modern method for importing external CMDB data when you want IRE processing with low custom technical debt?

A. IntegrationHub ETL / supported Service Graph ingestion pattern

B. Direct table writes

C. Arbitrary Business Rules

D. Manual inserts only

Why: Supported ingestion frameworks are designed to integrate with IRE and reduce custom transform logic.

Q30. An Asset Manager wants Asset and CI records to stay synchronized automatically. What two conditions are important?

A. One-to-one physical mapping between Asset and CI

B. Relevant Asset↔️CI synchronization business rules are active

C. A nightly clone

D. Manual updates only

Why: Asset/CI synchronization depends on a valid paired relationship and active synchronization logic.

Q31. How should Assigned To be synchronized between a CI and its associated Asset without building custom scripting?

A. Asset-CI Field Mapping

B. New custom field

C. Client Script

D. UI Action

Why: Asset-CI Field Mapping is the supported way to define field synchronization rules.

Q32. What relationship best describes an Application running on a Server?

A. Application → Runs On → Server

B. Application → Used by → Server

C. Application → Contains → Server

D. Application → Owns → Server

Why: Runs On is the standard semantic relationship for an application hosted on a server.

Q33. Which ServiceNow solutions automatically discover/create CI relationships from technical topology?

A. Discovery and Service Mapping

B. Workflow Studio and IntegrationHub ETL only

C. Change Management and Knowledge

D. HRSD and CSM

Why: Discovery and Service Mapping identify infrastructure/application dependencies and create relationships.

Q34. A Change Request must represent work across a group of CIs defined by a Dynamic CI Group. Which field is used to reference that group in the study scenario?

A. Configuration Item

B. Business Service

C. Service Offering

Why: Dynamic CI Groups can be referenced as configuration items for operational use cases involving groups of CIs.

Q35. Which service mapping approach is best known for precise, pattern-driven mapping of mission-critical application services?

A. Dynamic CI Group

B. Service Mapping (Top-down)

C. Tag-based only

D. Manual spreadsheet

Why: Top-down Service Mapping uses entry points and patterns for precise dependency mapping.

Q36. Which application-service approach is particularly suitable for cloud-native, container, or VM environments when resource tags are available?

A. Tag-Based

B. Dynamic CI Group only

C. Manual map

D. Data Certification

Why: Tag-based services leverage cloud/resource tags to define membership.

Q37. What is a strong use case for a Dynamic CI Group?

A. Grouping CIs dynamically with CMDB queries/filters for operational service contexts

B. Replacing all Business Applications

C. Running Discovery patterns

D. Storing financial transactions

Why: Dynamic CI Groups are query-driven CI collections useful when full dependency mapping is not required.

Q38. Which tool is appropriate for pre-built Store integrations that ingest third-party source data into the CMDB?

A. Service Graph Connector

B. ServiceNow Discovery only

C. Client Script

D. UI Builder

Why: Service Graph Connectors are packaged integrations for specific external data sources.

Q39. Which tool is an automated agentless discovery solution that runs probes/patterns against infrastructure?

A. Service Graph Connector

B. ServiceNow Discovery

C. CMDB Health

D. Data Certification

Why: Discovery actively identifies infrastructure and populates technical CI/relationship data.

Q40. A Service Graph Connector's transform definition is customized. If an upgrade later changes the baseline definition, what upgrade behavior should be expected?

A. The platform may create a skipped change rather than overwrite the customization

B. Upgrade always refuses to start

C. Upgrade deletes the connector

Why: Customized application records are generally protected from direct overwrite, leaving a skipped-change review path.

Q41. What is the support implication of customizing Service Graph Connector mappings?

A. Custom mapping changes become customer-maintained customizations

B. ServiceNow guarantees every customization

C. The connector cannot run again

D. No impact

Why: Customizations can increase maintenance responsibility and upgrade review effort.

Q42. Where is the Service Graph Connector Central experience commonly accessed?

A. Discovery Admin Workspace

B. CMDB Data Manager only

C. Service Operations Workspace only

Why: Connector administration is surfaced in Discovery Admin Workspace.

Q43. What should be used to maintain non-discoverable attributes such as warranty expiration dates when an authoritative external source owns the data?

A. Controlled import/integration from the authoritative source, processed through supported reconciliation

B. Manual edits only

C. A new CI class for each attribute

D. Disable reconciliation

Why: Non-discoverable business attributes should come from their authoritative system through governed integration and source precedence.

Q44. Which two scenarios are especially suitable for Agent Client Collector (ACC) compared with relying only on scheduled IP-based discovery?

A. Endpoints in environments where an installed agent is preferable for visibility

B. Endpoints such as laptops that may connect intermittently

C. Every Layer-2 network device must use ACC

D. ACC is only for data-center servers

Why: ACC supports push-based endpoint visibility and is useful for endpoint-style devices, including systems that are not continuously reachable by scheduled IP discovery.

Q45. On an Incident, how can CI data such as Support Group be used after a Configuration Item is selected?

A. It can provide routing/support context for the incident

B. It automatically changes the CI class

C. It disables Assignment Rules

D. It creates a new service offering

Why: ServiceNow's CSDM guidance states that CI data such as Support Group can be used to identify routing details and service impact.

Q46. What is the primary purpose of a Service Graph Connector?

A. Import and integrate third-party data into CMDB and, where supported, non-CMDB tables

B. Replace all identification rules

C. Create catalog items

D. Run client-side scripts

Why: ServiceNow describes Service Graph Connectors as packaged integrations for importing third-party data into the platform/CMDB.

Q47. Which pairing is correct?

A. ACC = agent-based collection; Discovery = agentless pattern-based discovery

B. ACC = agentless; Discovery = agent-based

C. Both require agents on every endpoint

D. Neither can populate the CMDB

Why: ACC places an agent on endpoints; ServiceNow Discovery is generally agentless and uses MID Server/pattern-based discovery.

Q48. For Microsoft SCCM data, which option normally provides the fastest time-to-value when a supported prebuilt connector is available?

A. Custom Import Set

B. Agent Client Collector

C. Service Graph Connector

D. Custom Business Rule

Why: A supported Service Graph Connector provides a packaged, upgrade-aware integration and usually reduces custom mapping work.

Q49. An authoritative asset system owns warranty-expiration data that Discovery cannot collect. What is the best maintenance pattern?

A. Create a new CI class for warranty dates

B. Load/integrate the attribute from the authoritative source and let IRE/reconciliation govern updates

C. Ask Discovery to invent the value

D. Store it only in work notes

Why: Non-discoverable attributes should come from their system of record through a governed integration.

Q50. Which supported feature should be preferred over a custom Business Rule to synchronize Assigned to between a CI and its paired Asset?

A. Asset-CI Field Mapping

B. UI Policy

C. Dictionary Override

D. Client Script

Why: Asset-CI Field Mapping is the purpose-built mechanism for paired Asset/CI field synchronization.

Q51. Which CMDB Health scorecard includes Duplicate, Orphan, and Stale CI metrics?

A. Correctness

B. Compliance

C. Completeness

Why: Duplicate, orphan, and stale CI checks are Correctness metrics.

Q52. Which two are examples of CMDB Health data-quality areas?

A. Duplicate CIs

B. Missing CIs as a distinct standard scorecard metric

C. Stale CIs

D. Upgraded CIs

E. Downgraded CIs

Why: Both are standard correctness concerns in CMDB Health.

Q53. What is required before CMDB Health scores are generated?

A. Nothing; all scores are calculated without jobs

B. Install an unrelated plugin

C. Activate/run the CMDB Health scheduled jobs

Why: CMDB Health metrics depend on scheduled health jobs to calculate results.

Q54. Match the CMDB Health areas: Required fields, Audit, Orphan.

A. Completeness → Required

B. Compliance → Audit

C. Correctness → Orphan

Why: Required fields measure completeness, audits measure compliance, and orphan checks are correctness.

Q55. A Server class has 60,000 CIs in Duplicate evaluation, but only 50,000 in Orphan evaluation. What most plausibly explains the reduced scope?

A. An inclusion rule narrows the Orphan metric population

B. Duplicate always ignores 10,000 CIs

C. A dashboard rendering defect

D. Reconciliation priority

Why: Health inclusion rules can scope which CIs are evaluated for a specific metric.

Q56. The Duplicate metric shows Healthy/Evaluated = 59,000/60,000. What does 60,000 represent?

A. The total CIs in scope for that metric

B. Only Linux servers

C. Only unhealthy CIs

D. The number of duplicate pairs

Why: Evaluated is the total population assessed by the metric.

Q57. A CMDB Data Foundations metric for Unique Locations is low. What is the safest remediation approach?

A. Delete one location without review

B. Keep both duplicates

C. Review duplicates, repoint dependent CIs to the valid location, then remove the duplicate

D. Treat the location as a CI duplicate

Why: Location duplicates should be resolved carefully so dependent records are not left pointing at the wrong record.

Q58. A CMDB administrator sees several poor Data Foundations metrics. What are two good next actions?

A. Focus on Critical/High priority metrics

B. Add exclusions only to make scores look better

C. Use linked remediation playbooks

D. Delete all retired CIs regardless of policy

Why: Prioritization and remediation guidance improve the underlying data rather than masking it.

Q59. Where would an administrator generally work with Natural Language Queries for CMDB analysis?

A. CMDB Health Dashboard

B. CMDB Workspace

C. CMDB Data Manager only

D. CI Class Manager only

Why: Natural-language CMDB experiences are surfaced in CMDB Workspace.

Q60. Which structured guidance should be used when a Data Foundations metric identifies an issue such as missing or duplicate foundational data?

A. The remediation/Get Well playbook linked to the relevant Data Foundations metric

B. A random fix script

C. Clone the instance

D. Rebuild the class hierarchy

Why: Data Foundations dashboards provide playbooks to guide remediation for the metric being reported.

Q61. Where in CMDB Workspace can Data Manager governance policies be created and managed?

A. Service Operations Workspace

B. Management area of CMDB Workspace

C. CI Class Manager only

D. CMDB 360 tab only

Why: CMDB Data Manager capabilities are exposed under Management in CMDB Workspace.

Q62. Where do users review tasks generated by CMDB Data Manager policies?

A. CMDB Workspace → My Work

B. CMDB Health Dashboard → Audit

C. Application Navigator → My Tasks only

D. CI Class Manager

Why: Policy tasks are surfaced to assignees in the My Work experience.

Q63. A CMDB administrator wants to remove Linux Servers that have not been updated for six months in a controlled, policy-driven way. Which Data Manager policy is most appropriate?

A. Archive policy

B. Business Rule

C. UI Policy

D. Client Script

Why: Archive policies provide governed lifecycle handling instead of ad hoc scripting.

Q64. What is the key difference between Attestation and Certification?

A. Attestation confirms the CI still exists/is valid; Certification validates specified CI attributes

B. They are identical

C. Certification is only for users

D. Attestation cannot be scheduled

Why: Attestation is existence/ownership-style confirmation; certification is attribute-level data validation.

Q65. Which policy type creates tasks for validating specified CI field values?

A. Retire

B. Attestation

C. Delete

D. Certification

Why: Certification policies are used for policy-driven validation of CI attributes.

Q66. Which policy type is used to automate tasks that ask responsible users to confirm that CIs still exist or are still valid?

A. Retire

B. Attestation

C. Delete

D. Certification

Why: Attestation tasks are designed for acknowledgement/validation of the CI's continued relevance or existence.

Q67. A team is transitioning from legacy CMDB status fields to CSDM lifecycle fields. Which mapping table is used to map legacy values?

A. life_cycle_stage

B. life_cycle_stage_status

C. life_cycle_control

D. life_cycle_mapping

Why: The lifecycle mapping table connects legacy status values with CSDM lifecycle stages/statuses.

Q68. Match the lifecycle objects: life_cycle_stage, life_cycle_stage_status, life_cycle_mapping, life_cycle_object.

A. Stage → meta-level lifecycle state

B. Stage Status → sub-level lifecycle state

C. Mapping → legacy-to-lifecycle mapping

D. Object → defines lifecycle applicability by object/CI type

Why: These records define lifecycle vocabulary, detailed status, mappings, and applicability.

Q69. What is the most critical factor for operationalizing Configuration Management successfully?

A. Populate as much data as possible

B. Use automation only

C. Let all teams freely change CMDB records

D. Establish governance and continuously monitor CMDB health

Why: Governance, ownership, source control, and ongoing health management are essential for sustainable CMDB quality.

Q70. A CMDB administrator wants server patch compliance to appear meaningfully on the Compliance scorecard. What must exist conceptually?

A. Certification/audit policy and filters plus scheduled evaluation

B. Only Duplicate and Orphan rules

C. Only a Dynamic CI Group

D. A clone

Why: Compliance requires defined expected values/audit logic and jobs that evaluate the scoped CIs.

Q71. What is the preferred workspace-based way to manage duplicate CIs?

A. Deduplication Task module only

B. My Tasks in navigator only

C. Deduplication Dashboard in CMDB Workspace

Why: The workspace dashboard provides a consolidated remediation experience.

Q72. Two duplicate CIs exist: the discovered CI has the latest IP address, while the manually imported CI has a valuable relationship to a critical Business Application. What is the best remediation idea?

A. Delete one CI without review

B. Keep the discovered CI as main and merge/carry over the valid relationship from the other CI

C. Keep both indefinitely

D. Copy all values from the manual CI regardless of freshness

Why: Duplicate remediation should preserve the best authoritative attributes and valid relationships while converging to one main CI.

Q73. Which workspace area contains deduplication tasks generated from IRE-related duplicate detection?

A. My Work / task status area in CMDB Workspace

B. Insights only

C. Home only

D. Knowledge

Why: Generated remediation tasks are surfaced through the workspace's work/task experience.

Q74. Which role is generally the highest-level CMDB administration role in the study role-matching set?

A. CI Analyst

B. CMDB Process Owner

C. Configuration Manager / CMDB Administrator

D. Service Owner

Why: This role carries the broadest CMDB administration privileges among the listed choices.

Q75. Which role is typically accountable for managing portfolio/service elements through their lifecycle?

A. CI Analyst

B. CMDB Process Owner

C. Service or Product Owner

D. End user

Why: Service/product owners own lifecycle and quality outcomes for their assigned services or products.

Q76. Which role commonly consumes CMDB reports and dashboards with read-focused access?

A. CI Analyst

B. CMDB Administrator

C. System Administrator

D. Application Developer

Why: CI analysts often work with CMDB reporting and analysis without the highest administrative rights.

Q77. What is true about the CSDM and CMDB Data Foundations dashboards in current base systems?

A. They require a separate paid application in every instance

B. Their capability is provided by com.snc.cmdb.getwell, which is activated by default in base systems

C. They are available only with Service Mapping Plus

D. They exist only in Performance Analytics

Why: ServiceNow documents the dashboards as part of the CMDB Get Well capability activated by default in base systems.

Q78. A CMDB Data Manager policy requires task review/approval by the group responsible for the target CIs. Which CI field is used for that group assignment?

A. Approval Group

B. Managed by Group

C. Owned by Group

D. Assignment Group

Why: CMDB Data Manager uses the Managed by Group attribute for policy-task review/approval assignment behavior.

Q79. Which Data Foundations metric specifically helps identify hardware CIs that may be prone to identification problems because serial numbers are missing?

A. Hardware CIs Missing Serial Numbers

B. Orphan Business Services

C. Duplicate User Groups

D. Stale Change Requests

Why: The CMDB Data Foundations dashboard explicitly includes a Hardware CIs Missing Serial Numbers metric.

Q80. Before Archive or Delete life-cycle policies target a CI, what state is normally required by CMDB Data Manager?

A. The CI must be retired according to its retirement definition

B. The CI must be newly discovered

C. The CI must have no sys_id

D. The CI must be a Principal Class

Why: Archive and Delete life-cycle policies enforce retirement definitions and the appropriate retired state.

Q81. What is the base-system effective duration commonly used by the default CMDB Health staleness rule?

A. 7 days

B. 30 days

C. 60 days

D. 365 days

Why: ServiceNow documents a default 60-day effective duration, which can be overridden for specific classes.

Q82. Which field is evaluated by CMDB Health staleness rules to determine whether a CI has been updated within the effective duration?

A. sys_created_on

B. sys_updated_on

C. first_discovered

D. discovery_source

Why: Staleness is based on whether the CI has been updated within the rule effective duration.

Q83. For a duplicate set containing three CIs, how many are counted as duplicate CIs by the CMDB Health Duplicate metric?

A. 1

B. 2

C. 3

D. 0

Why: For a duplicate set, ServiceNow counts total CIs in the set minus one.

Q84. What does a Health Inclusion Rule do?

A. Narrows which CIs/classes are included in a CMDB Health calculation

B. Decides which discovery source wins reconciliation

C. Creates a new class

D. Reclassifies a CI

Why: Health Inclusion Rules control evaluation scope; they do not replace IRE rules.

Q85. Which policy families should you recognize for current CMDB Data Manager study purposes?

A. Retire, Archive, Delete, Attestation, Certification

B. Incident, Problem, Change, Request, Knowledge

C. Duplicate, Switch, Upgrade, Downgrade, Merge

D. Discovery, ACC, ETL, SGC, Import Set

Why: Current CMDB Data Manager includes life-cycle policy types plus Attestation and Certification experiences.

Q86. What is the simplest way to remember Attestation versus Certification?

A. Attestation = confirm the CI exists/is still valid; Certification = validate specified data values

B. Attestation = delete; Certification = archive

C. Both only find duplicates

D. Certification is only for users

Why: Attestation confirms continued existence/validity; Certification checks defined fields and values.

Q87. Where is Unified Map accessed?

A. CMDB Workspace

B. CI Class Manager

C. CMDB Data Manager only

Why: Unified Map is a CMDB Workspace capability.

Q88. On Unified Map, where can a user view contextual related records such as active incidents or planned changes for the home node?

A. Map canvas only

B. Content controls

C. Contextual side panel

D. Toolbar

Why: The contextual side panel surfaces related operational records without leaving the map.

Q89. Which two filters are available for hiding irrelevant CI nodes in Unified Map?

A. Discovery source

B. CI class/type

C. Approval state

D. Knowledge base

Why: Unified Map supports filtering by CI class and discovery source, as well as relationship type, location, ownership, and layer count.

Q90. A CMDB query needs Seattle databases connected to application services and also related Incidents. What two actions are essential?

A. Filter the Database node for Location = Seattle

B. Add Incident as a non-CMDB table to the query

C. Add only property columns

D. Use a scheduled job

Why: Use a CMDB node filter for location and bring Incident into the query as a non-CMDB table.

Q91. Where can CMDB 360 saved queries be viewed/created in CMDB Workspace?

A. Saved Queries window on Insights

B. CMDB Query Builder only

C. Coverage window

D. Saved Queries window on the CMDB 360 tab

Why: CMDB 360 has a dedicated saved-query experience for multisource analysis.

Q92. Which are core business values of a healthy CMDB? Choose the strongest answers.

A. Streamlining Incident and Change Management

B. Strengthening operational resilience

C. Collecting arbitrary financial data as the primary purpose

D. Improving impact/context through trusted CI relationships

Why: A trusted CMDB supports operational processes, resilience, and impact analysis. Financial management may integrate with CMDB but is not its primary purpose.

Q93. Which group-assignment attributes can be synchronized from a Technology Management Offering through its Dynamic CI Group to member CIs?

A. Support Group

B. Change Group

C. Managed by Group

D. Owned by Group

Why: ServiceNow CSDM documentation names Support Group, Change Group, and Managed by Group as the synchronized group-assignment attributes.

Q94. For a Certification policy, what does enabling “Allow field updates” permit the reviewer to do?

A. Change certified field values while reviewing the CI

B. Modify identification rules

C. Reclassify the CI automatically

D. Disable IRE

Why: Certification reviewers can be allowed to correct/update field values as part of the certification task.

Q95. Which Unified Map filters are definitely supported for filtering CI nodes?

A. CI class and discovery source

B. Knowledge base and catalog category

C. Approval state and SLA

D. Survey metric and HR profile

Why: Unified Map filtering includes CI class, discovery source, relationship type, location, ownership, and layer count.

Q96. What does the right-hand Evaluated value in a CMDB Health Healthy/Evaluated display represent?

A. Total CIs in scope for that metric

B. Only unhealthy CIs

C. Duplicate pairs only

D. Newly created CIs only

Why: Evaluated is the population assessed by the metric; Healthy is the subset that passed.

Q97. When a Data Foundations metric is poor, what is the purpose of its playbook?

A. Provide structured analysis and remediation guidance

B. Automatically rewrite all CMDB data without review

C. Replace IRE

D. Disable the metric

Why: Playbooks provide structured problem context and remediation guidance.

Q98. Which CMDB feature is best for visually exploring CI relationships and related operational context?

A. Unified Map

B. Update Set Preview

C. ACL Debugger

D. Knowledge Search

Why: Unified Map provides a relationship-oriented visual view and related context.

Q99. A CMDB administrator wants to improve data quality specifically in relation to CSDM adoption. Which dashboard should be used?

A. Default CMDB Health Dashboard

B. CSDM Data Foundations Dashboard

C. ServiceNow Health Scan

Why: The CSDM Data Foundations Dashboard is designed to highlight data gaps and adoption issues aligned to CSDM stages.

Q100. What is the prescribed CSDM rollout sequence?

A. Foundation, Crawl, Walk, Run, Fly

B. Architecture, Business, Technical, Governance

C. Initiate, Plan, Execute, Deliver, Close

D. Initial, Developing, Defined, Managed

Why: CSDM adoption is commonly described through the Foundation, Crawl, Walk, Run, and Fly progression.

Q101. A customer is aligned to the CSDM Walk stage. Which benefit is most directly enabled?

A. Impact assessment for Incident, Problem, and Change on Business Services

B. APM rationalization velocity only

C. Additional OLA stratification only

Why: At Walk, service context and relationships support better operational impact analysis.

Q102. Which CSDM domain explains relationships among infrastructure CIs, Technical Service Offerings, and Application Services?

A. Service Delivery

B. Design & Planning

C. Build & Integration

D. Foundation

E. Service Consumption

Why: The Service Delivery / Manage Technical domain focuses on the technical services and supporting operational infrastructure.

Q103. Where is regulatory or compliance-oriented information modeled in CSDM?

A. Customer Service Offerings and Databases

B. Business Applications and Information Objects

C. Technology Management Service Offerings and Dynamic CI Groups

Why: Information Objects describe the logical information used or managed by Business Applications and can support governance and regulatory context.

Q104. What is the purpose of an Information Object?

A. Describes data in general on a group of CIs

B. Describes logical data associated with Business Applications

C. Describes data exchanged only between an API interface and an application

Why: Information Objects model important logical information used by business applications.

Q105. Which CSDM objects best represent the Sell / Consume domain?

A. Application Service, Technical Service, Technical Service Offering

B. Business Service and Business Service Offering

C. Business Capability, Information Object, Business Application

Why: Sell/Consume is centered on customer-facing business services and their offerings.

Q106. A CMDB Manager is adding Design & Planning domain components. Which stakeholders are especially important?

A. Enterprise Architect

B. Application Service Owner

C. Business Relationship Manager

D. Application Owner

Why: Design & Planning focuses on business applications, capabilities, processes, and architecture ownership.

Q107. Which CSDM-related data is especially valuable during a critical incident to understand affected configuration items?

A. Application Service environment attribute

B. Incident history of similar CIs

C. Affected CI (task_ci) related list

D. Service Offerings by Department or Location

Why: The Affected CIs list captures additional CIs impacted by the task and supports impact assessment.

Q108. How does Change Management benefit from CSDM? Choose two.

A. Identify blackout windows for the change

B. Understand impact of the change on services

C. Route the change dynamically by itself

D. Determine root cause of the change issue

Why: Service and offering context can inform schedules/blackouts and provide downstream business/technical impact.

Q109. How can Application Service Mapping help Change Management?

A. Understand the physical location of CIs only

B. Understand business and service impact of affected CIs

C. Identify which device goes offline first

Why: A service map shows dependencies so change impact can be evaluated more accurately.

Q110. Which stage commonly introduces Technology Management Service Offerings that can provide support/ownership context to related CIs?

A. Crawl

B. Walk

C. Foundation

D. Fly

E. Run

Why: Walk expands operational service modeling and ownership context.

Q111. Where do Business Application and Information Object belong in the current CSDM conceptual model?

A. Design & Planning

B. Service Consumption

C. Service Delivery

D. Foundation only

Why: ServiceNow documents Business Application, Business Capability, and Information Object in Design & Planning.

Q112. Which CSDM object represents sensitive information types such as PII, PCI, or HIPAA-related data used by a Business Application?

A. Information Object

B. Dynamic CI Group

C. Hardware Model

D. Service Offering

Why: Information Objects represent logical information used by Business Applications and can identify sensitive/regulatory data.

Q113. What is the easiest way to avoid mixing up CSDM domains with implementation stages?

A. Domains describe model purpose; stages describe adoption maturity: Foundation → Crawl → Walk → Run → Fly

B. They are the same thing

C. Stages are CI classes

D. Domains are dashboard scores

Why: Domains organize the conceptual model; stages describe progressive adoption.

Q114. A CI class must be designated as a Principal Class. Where is this configured?

A. CI Class Manager > Class Info / Basic Info

B. CMDB Health > Correctness

C. CMDB Data Manager > Policies

D. Discovery Dashboard

Why: Principal Class designation is class metadata managed from CI Class Manager, not a health or lifecycle policy.

Q115. A lower-priority source should be allowed to update a CI only after the higher-priority source has been inactive for a defined time. What should be configured?

A. Data Refresh Rule

B. Health Inclusion Rule

C. Identification Rule

D. De-duplication Template

Why: A Data Refresh Rule works with reconciliation precedence to permit a lower-priority source after the configured inactivity period.

Q116. When a CI is reclassified from one sibling class to another sibling class, how is the operation categorized?

A. Upgrade

B. Downgrade

C. Switch

D. Retirement

Why: A sibling-to-sibling reclassification is a class switch.

Q117. You need to ingest Microsoft SCCM data and a supported packaged connector is available. Which option is normally preferred?

A. Custom Import Set

B. Agent Client Collector

C. Service Graph Connector

D. Client Script

Why: A supported Service Graph Connector is the packaged integration path and generally reduces custom mapping and maintenance.

Q118. Which statement best describes the purpose of Service Graph Connector field mappings?

A. They map source data into the target ServiceNow/CMDB model

B. They replace identification rules

C. They calculate CMDB Health

D. They create ACLs

Why: Connector mappings translate source-system fields into the ServiceNow data model; IRE and reconciliation still govern CI identity and authority.

Q119. An Asset and its paired CI must stay synchronized for mapped fields. Which native mechanism should be considered before custom scripting?

A. Asset-CI Field Mapping / synchronization

B. UI Policy

C. Dictionary Override

D. Knowledge workflow

Why: Use the platform's Asset-CI synchronization capability before introducing custom logic.

Q120. Which two outcomes are major reasons CMDB context improves Event Management?

A. Better root-cause/impact context

B. Improved understanding of services at risk

C. Automatic deletion of every alert

D. Removal of all monitoring tools

Why: Event Management uses CMDB/service context to correlate technical signals, identify likely root cause, and understand affected services.

Q121. Which mapping approach groups CIs using common criteria rather than discovering a dependency path from an entry point?

A. Dynamic CI Group

B. Top-down Service Mapping only

C. Class downgrade

D. Certification

Why: Dynamic CI Groups collect CIs that match defined criteria and can be used as group CIs in operational processes.

Q122. Which Service Mapping approach uses metadata labels from environments such as cloud/container platforms to build application-service membership?

A. Tag-based mapping

B. Certification

C. Data Manager archive

D. CMDB Health audit

Why: Tag-based mapping uses tags/labels to identify and group relevant resources into application services.

Q123. A CI fails the Recommended-fields health test and the team wants a task created automatically. What setting must be enabled for that metric?

A. Create Task

B. Independent

C. Principal Class

D. Allow fallback

Why: CMDB Health metric preferences include a Create Task option and a Task Assignee Group for failed metric tests.

Q124. A remediation workflow for a failed CMDB Health test is selected through which configuration object?

A. CMDB Remediation Rule

B. Identification Entry

C. Health Inclusion Rule

D. Service Offering

Why: A CMDB Remediation Rule references the remediation workflow and controls task type, filter, and manual/automatic execution.

Q125. Which two health metrics use specialized remediation rather than the normal one-task-per-failed-CI pattern?

A. Duplicate

B. Audit

C. Recommended

D. Stale

Why: Duplicate failures use de-duplication tasks, while audit failures use audit tasks; other health metrics can create standard health tasks.

Q126. For the base-system staleness rule, which field is evaluated to determine whether the CI has been updated recently enough?

A. Updated (sys_updated_on)

B. Created (sys_created_on)

C. Name

D. Discovery source only

Why: Staleness is based on how recently the CI was updated relative to the rule's effective duration.

Q127. Which statement about the CMDB Health Relationships KPI is correct?

A. It evaluates relationship quality such as orphan/duplicate relationships and relationship compliance

B. It is another name for Completeness

C. It only measures stale CIs

D. It replaces IRE

Why: Current CMDB Health also evaluates relationship health in addition to the core CI data-quality KPIs.

Q128. Which five items are CMDB Data Manager policy families rather than de-duplication or reclassification operations?

A. Retire

B. Archive

C. Delete

D. Attestation

E. Certification

F. De-duplication

G. Reclassification

Why: Retire, Archive, Delete, Attestation, and Certification are Data Manager policy experiences; de-duplication and reclassification are handled through other CMDB capabilities.

Q129. A CMDB Health metric shows 720 Healthy / 800 Evaluated. What does 800 represent?

A. Total CIs evaluated for that metric

B. Only failed CIs

C. Only duplicate CIs

D. Only newly discovered CIs

Why: Evaluated is the population tested by the metric; Healthy is the subset that passed.

Q130. A Business Application handles regulated patient information. Which CSDM object is appropriate for representing the logical information being used?

A. Information Object

B. Dynamic CI Group

C. Hardware Model

D. Change Model

Why: Information Objects represent logical information used by Business Applications and can model sensitive or regulated information types.

Q131. Which statement correctly separates a Business Application from an Application Service / service instance?

A. Business Application is a design/planning view; Application Service represents an operational deployment/service instance

B. They are always the same CI

C. Business Application belongs only to Service Consumption

D. Application Service is only a financial record

Why: CSDM separates design-time application portfolio concepts from operational service instances used to deliver and manage technology.

Q132. A healthcare provider faces a critical incident affecting its patient management system. Which CSDM-related data is most useful for seeing all CIs affected by the incident?

A. Application Service environment attribute

B. Incident history of similar CIs

C. Affected CI (task_ci) related list

D. Service Offerings by Department or Location

Why: The Affected CI related list gives the incident/change context for the CIs associated with the task.

Q133. A CMDB Administrator wants to create a new CI class for an IoT Sensor. Which two practices best align with the class-model approach? (Choose 2)

A. Delete an unused class and replace it

B. Install/update the CMDB CI Class Models Store application and first verify that a suitable class does not already exist

C. Add the new class under the appropriate parent class

D. Modify an unrelated existing class

Why: Reuse the supported class model where possible; if a class is genuinely needed, extend the correct CI parent rather than repurposing unrelated classes.

Q134. A Service Graph Connector script transform was customized. During a later upgrade, ServiceNow ships an update to that same default definition. What should the administrator expect?

A. A skipped change is created and the local customized definition is not silently overwritten

B. The connector upgrade refuses to start

C. The upgrade stops with a fatal error

Why: A local customization can cause the incoming update to be treated as a skipped change rather than overwriting the customized record.

Q135. Which CI Class Manager feature helps a class owner quickly return to the Windows Server class they manage frequently?

A. Pinned Classes

B. Search CI Classes only

C. CI Favorites

Why: Pinned Classes provide quick access to classes a user works with frequently.

Q136. A Platform Owner is building CSDM governance. Match the governance scope to the stakeholders most directly involved.

A. Design domain -> Enterprise Architects + Platform Owner

B. Foundation domain -> Enterprise Architects + Data Stewards + Process Owners + Platform Owner

C. Design domain -> End users only

D. Foundation domain -> Application developers only

Why: CSDM governance is cross-functional: design needs architecture/platform ownership, while foundation data also needs stewardship and process ownership.

Q137. A CMDB Administrator is preparing CI data for an Integrated Risk Management implementation. Which two CSDM relationships provide useful application and information context? (Choose 2)

A. Logical CIs related directly to Information Objects

B. Business Applications related to Information Objects

C. Locations with parent records only

D. Business Applications related to Application Services

Why: Business Applications connect portfolio/design context to both logical information and operational application-service context.

Q138. The Duplicate metric shows Healthy/Evaluated = 59,000/60,000 for the Server class. What does 60,000 represent?

A. The total server CIs evaluated for that metric

B. Only Linux servers

C. Only unhealthy records

D. The number of duplicate sets

Why: Evaluated is the population tested by the metric; Healthy is the subset that passed.

Q139. A CMDB CI Class Owner wants to change the icon/label presentation for the UNIX Server class. Which CI Class Manager area is used?

A. Suggested Relationships

B. CI List

C. Attributes

D. Label Job

Why: The class label/icon presentation is handled through the class-label configuration experience rather than attribute or relationship configuration.

Q140. An organization is adding and removing classes from the Principal Class Filter. Which two administrative roles are required in the study scenario? (Choose 2)

A. sn_custom_admin

B. sn_cmdb_admin

C. cmdb_query_builder

D. personalize_dictionary

Why: The scenario is about administration of the CMDB principal-class configuration, not query-building or dictionary personalization.

Q141. The Change Management team wants one Change to target a Dynamic CI Group representing multiple CIs. Which Change Request field is populated with the Dynamic CI Group?

A. Configuration Item

B. Business Service

C. Service Offering

Why: Dynamic CI Groups can be selected through the Configuration Item context to represent a governed group of CIs.

Q142. A company asks which use case specifically needs Information Objects in its CSDM model. Which option is the strongest fit?

A. Customer Service wants proactive case management that needs application/data context

B. Asset Management only wants hardware lifecycle compliance

C. Event Operations only wants events converted to incidents

D. A team only wants device discovery

Why: Information Objects represent logical information used by Business Applications and become relevant when a use case needs that information/data context.

Q143. Which statement best describes the relationship between an Application and the Server on which it executes?

A. Application -> Runs On -> Server

B. Application -> Used by -> Server

C. Application -> Owns -> Server

D. Application -> Approves -> Server

Why: The Runs On relationship expresses that the application executes on the supporting server.

Q144. A CMDB Administrator wants structured guidance for missing serial numbers that may lead to duplicate Hardware CIs. What should they use?

A. CSDM Now Create playbooks

B. CMDB Health/Data Foundations remediation playbook guidance for the affected metric

C. A custom Business Rule as the first step

D. A Change Model

Why: Use the dashboard's structured remediation/playbook guidance to investigate the data-quality issue instead of starting with ad-hoc customization.

Q145. A Change Manager wants CI ownership information to support Change processing. If the selected CI is governed by its managing group, which base CMDB group field represents that ownership context?

A. Assignment group

B. Change group

C. Approval group

D. Managed by group

Why: Managed by Group identifies the group responsible for managing the CI; do not confuse it with task Assignment Group.

Q146. What is the purpose of an Information Object in CSDM?

A. Describe arbitrary text on a group of CIs

B. Represent logical information used by Business Applications

C. Represent only a physical database server

D. Replace an Application Service

Why: An Information Object models logical data/information such as patient, payment, customer, or employee information used by Business Applications.

Q147. Drag and Drop - Match each CSDM object to the most appropriate domain.

Item	Correct match
Business Process	**Design & Planning**
Business Application	**Design & Planning**
Application Service / Service Instance	**Service Delivery**
Business Service	**Service Consumption**

Why: Business Process and Business Application sit in Design & Planning; the deployed Application Service belongs to Service Delivery; Business Service belongs to Service Consumption.

Q148. Drag and Drop - Match each CMDB Health metric to the best description.

Item	Correct match
Duplicate	**Same real CI represented more than once**
Required	**Mandatory/required field is missing**
Orphan	**Expected CI relationships are missing**
Stale	**CI has not been updated within the staleness window**
Recommended	**Best-practice field is missing**
Audit	**CI values are evaluated against an expected/defined state**

Why: Keep the scorecard families straight: Required/Recommended -> Completeness; Duplicate/Orphan/Stale -> Correctness; Audit -> Compliance.

Q149. Drag and Drop - Match the ingestion architecture to the ServiceNow tool.

Item	Correct match
Agent-based automated collection	**Agent Client Collector (ACC)**
Organization-built load using staging and transform maps	**Import Sets / Transform Maps**
Packaged third-party vendor integration	**Service Graph Connector**
Agentless infrastructure discovery using patterns	**ServiceNow Discovery**

Why: The quickest clue is agent-based = ACC, agentless patterns = Discovery, packaged vendor integration = Service Graph Connector.

Q150. Drag and Drop - Match the service type to its meaning.

Item	Correct match
Application Service / Service Instance	**Operational representation of a deployed application/service**
Technology Management Service	**Technical service delivered and managed by technology teams**
Business Service	**Business-facing service consumed by business users/customers**

Why: Separate operational application topology, technical delivery, and business-facing consumption.

Q151. Drag and Drop - Match the Application Service creation/mapping approach to the best-fit scenario.

Item	Correct match
Top-down Service Mapping	**Mission-critical application needing precise pattern-driven dependency mapping**
Connection Suggestions	**Dynamic/cloud/container environment where observed connections help suggest membership**
Tag-based	**Resources consistently labeled with application tags**
Dynamic CI Group	**Small/simple service population defined by filters or saved CMDB queries**

Why: The exam may describe the environment rather than name the mapping method, so recognize the design pattern.

Q152. Drag and Drop - Match the CSDM object to the best-fit conceptual area.

Item	Correct match
Business Application	**Design & Planning**
Application Service / Service Instance	**Service Delivery**
Business Service	**Service Consumption**
Information Object	**Information portfolio / design and risk context for Business Applications**

Why: Do not confuse a Business Application (design/portfolio) with an Application Service (operational deployment).

Q153. Drag and Drop - Match the product/capability to the value CMDB context can provide.

Item	Correct match
Hardware Asset Management	**Asset/CI lifecycle and operational context for managed hardware**
Risk / Security Operations	**Criticality and infrastructure context for risk prioritization**
Discovery	**Detailed technical CI attributes and relationships**
Service Portfolio Management	**Service lifecycle and portfolio context connected to operational services**

Why: These mappings test the business value of a trusted CMDB across products, not just CMDB administration.

4. Final 48-Hour Revision Checklist

☐ Explain Identification vs Reconciliation without notes.

☐ Map Required/Recommended, Duplicate/Orphan/Stale, and Audits to the right Health KPI.

☐ Choose between Discovery, ACC, Service Graph Connector, IntegrationHub ETL, and Import Sets.

☐ Explain why CMDBTransformUtil belongs in an onBefore Transform Map script.

☐ Distinguish Attestation, Certification, Retire, Archive, and Delete.

☐ Explain CMDB 360 and why dynamic reconciliation needs it.

☐ Use Query Builder logic: nodes, relationships, node filters, non-CMDB tables.

☐ Distinguish CSDM domains from Foundation/Crawl/Walk/Run/Fly stages.

☐ Place Business Application, Information Object, Service Instance/Application Service, and Business Service correctly.

☐ Explain CMDB value for Incident, Change, Event Management, security/risk and resilience.

Exam-day mindset

- Before answering, check whether the question says Choose 2 / Choose 3 or requires drag-and-drop; select exactly the requested number.
- When a scenario clearly matches a supported ServiceNow capability, prefer that supported approach over unnecessary customization.
- Separate ingestion from governance: loading data is not the same as deciding source authority.
- Flag uncertain questions and move on instead of spending several minutes on one memory.

ONE LAST THING

You’ve done the prep. Now trust what you know.

GOOD LUCK

You’re closer than you think :)

Read carefully. Think in ServiceNow terms. Don’t overthink the obvious.